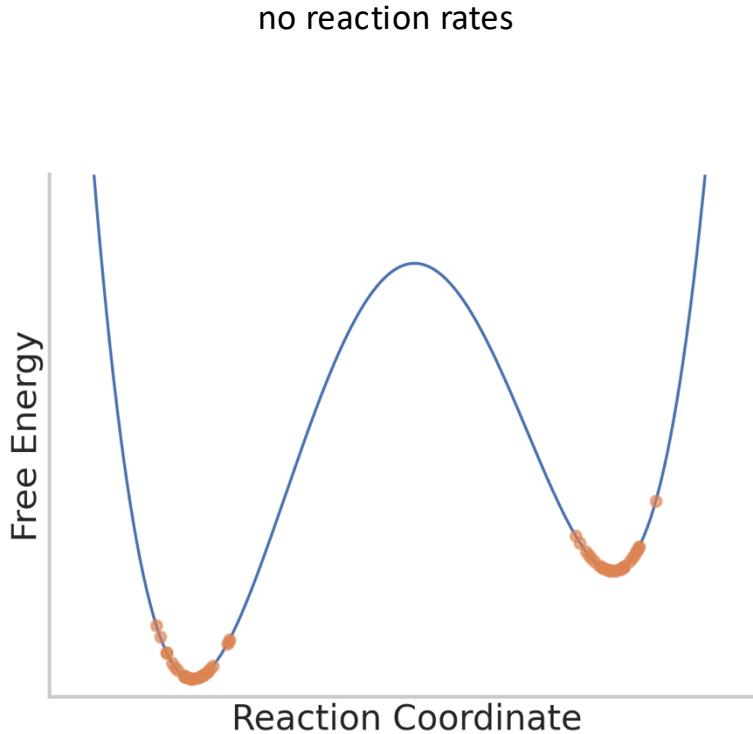


Finding Transition States with HIP

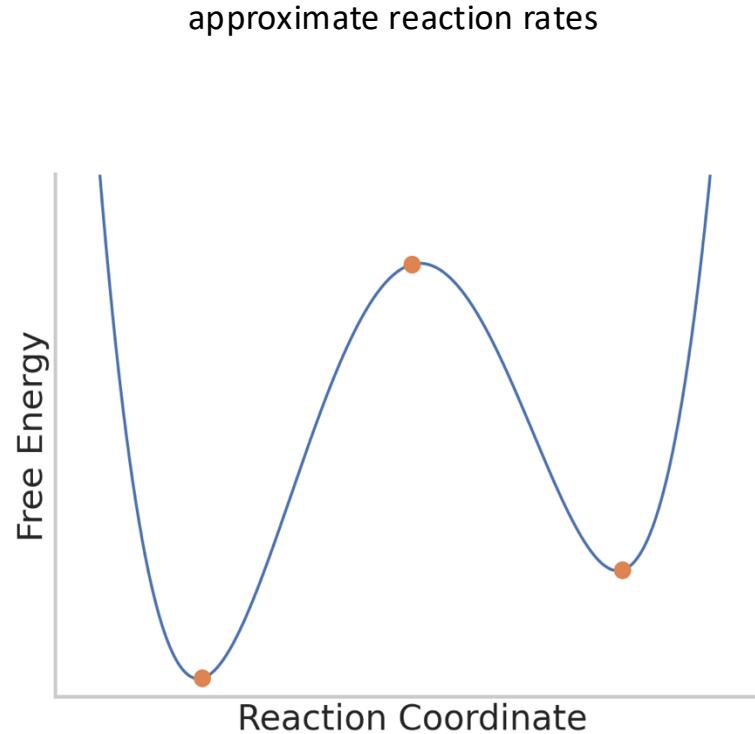
Gentles Ascent Dynamics (GAD)
and
Eigenvalue Descent
By Memo

With transition states we can approximate reaction rates

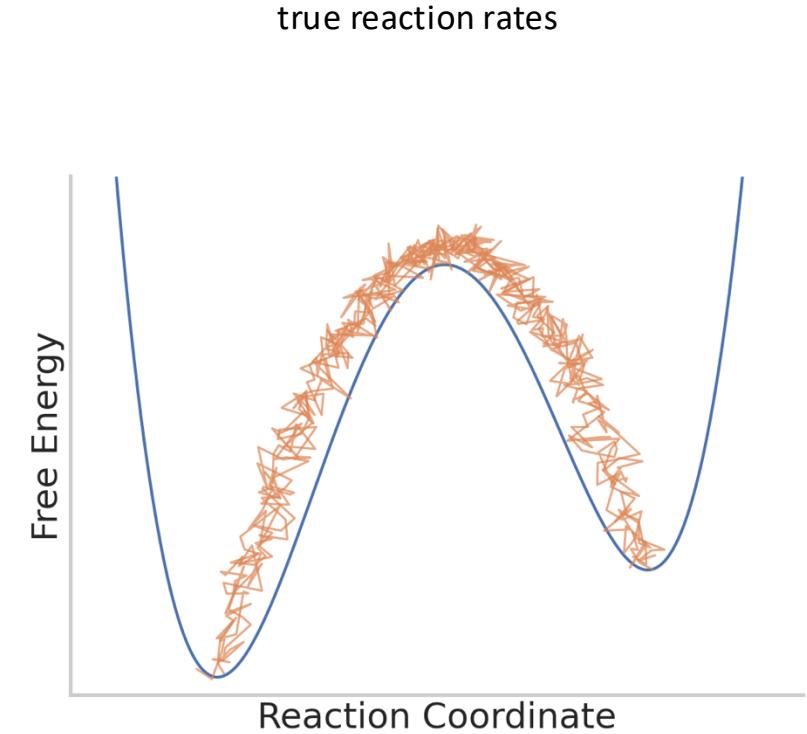
Free Energy Difference



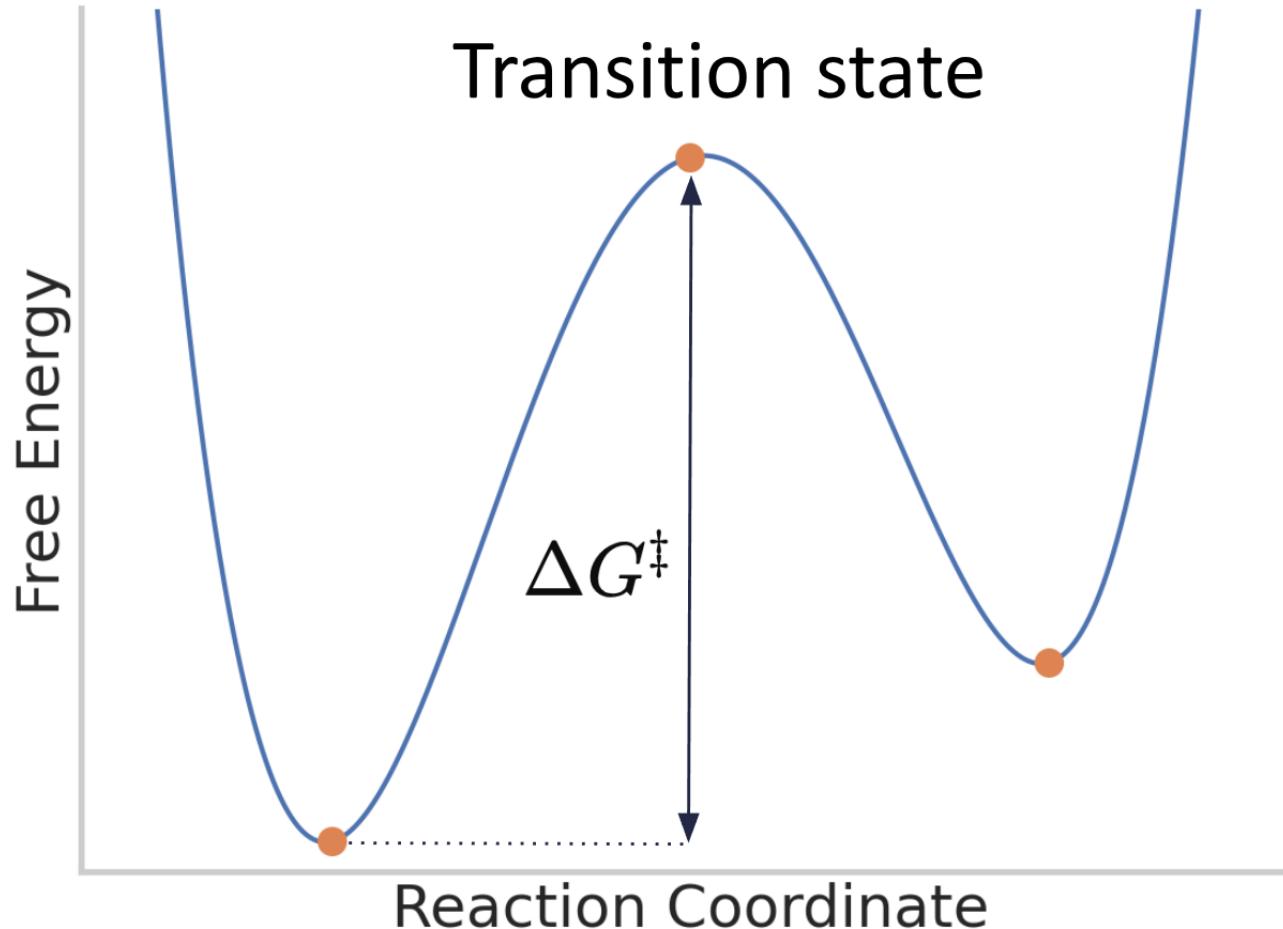
Transition State Theory



Transition Path Sampling



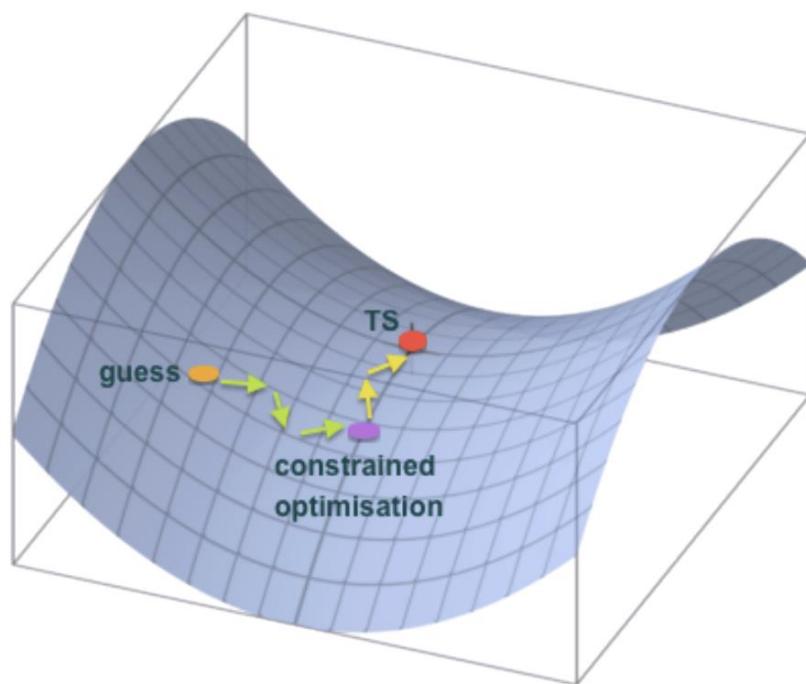
With transition states we can approximate reaction rates



Eyring equation:

$$k(T) = \frac{k_B T}{h} \exp\left(-\frac{\Delta G^\ddagger}{RT}\right)$$

We can identify transition states using the smallest two eigenvalues



Transition State = Saddle point in PES

Characterized by: $\lambda_0 < 0, \lambda_1 > 0, \lambda_2 > 0, \dots$

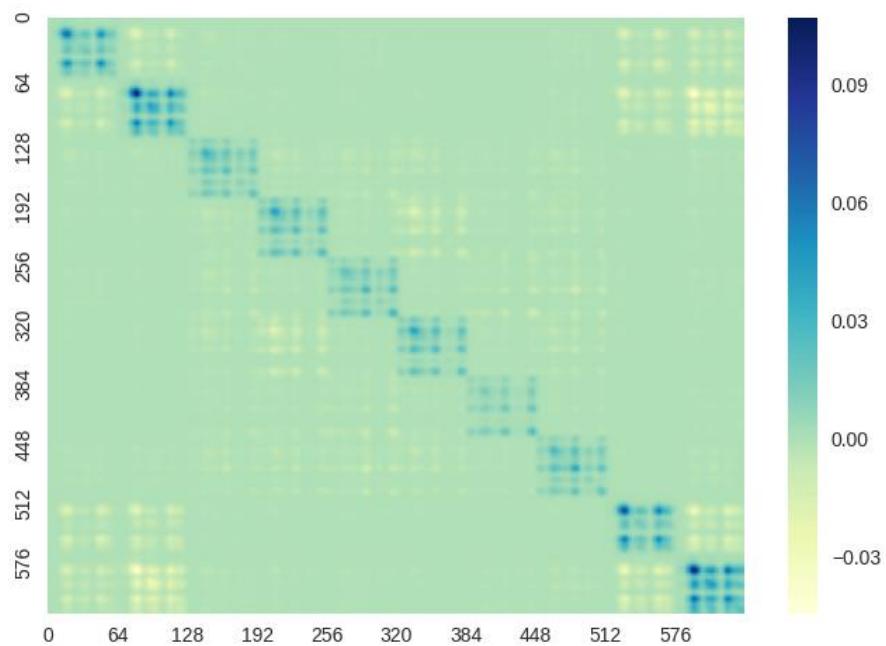
Product $\lambda_0 \times \lambda_1 < 0$ indicates TS

Eigvals, eigvecs = torch.eigh(H_f)

Eigvals = [λ₀, λ₁, ...] # λ₀, λ₁ are the smallest two eigenvalues

HIP predicts Hessians directly from atomic positions

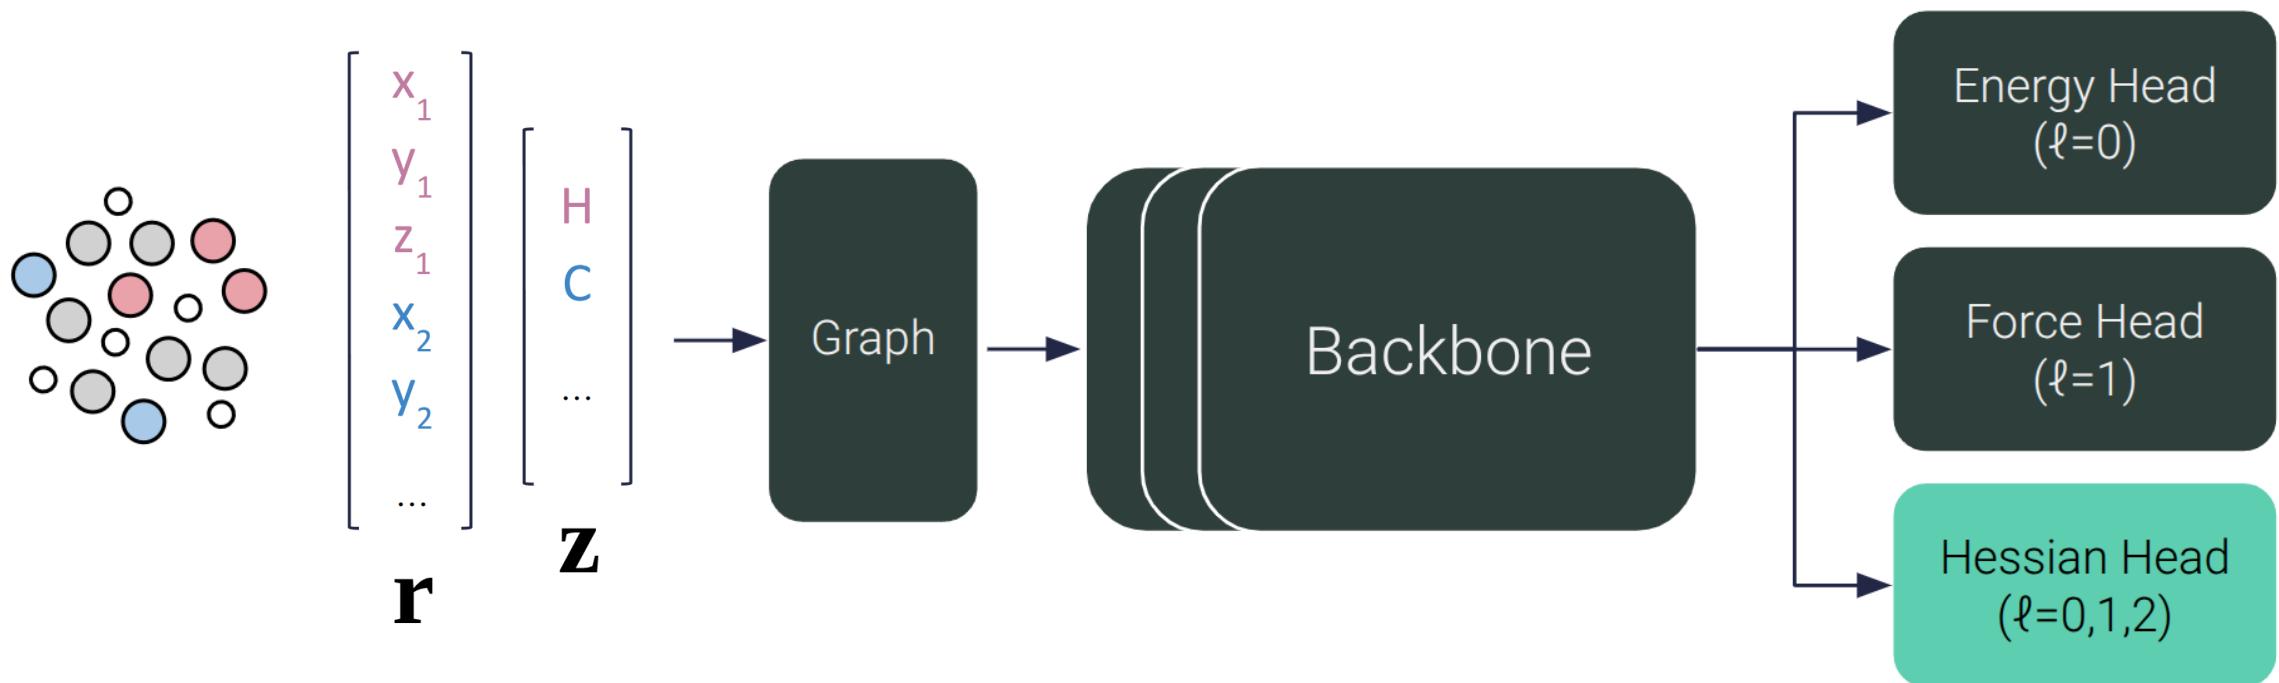
- HIP: Machine Learning Interatomic Potential
- Input: Atomic positions (x, y, z)
- Output: Full Hessians ($\nabla^2 V$)
- Key: No autodiff needed
- Enables direct saddle point analysis



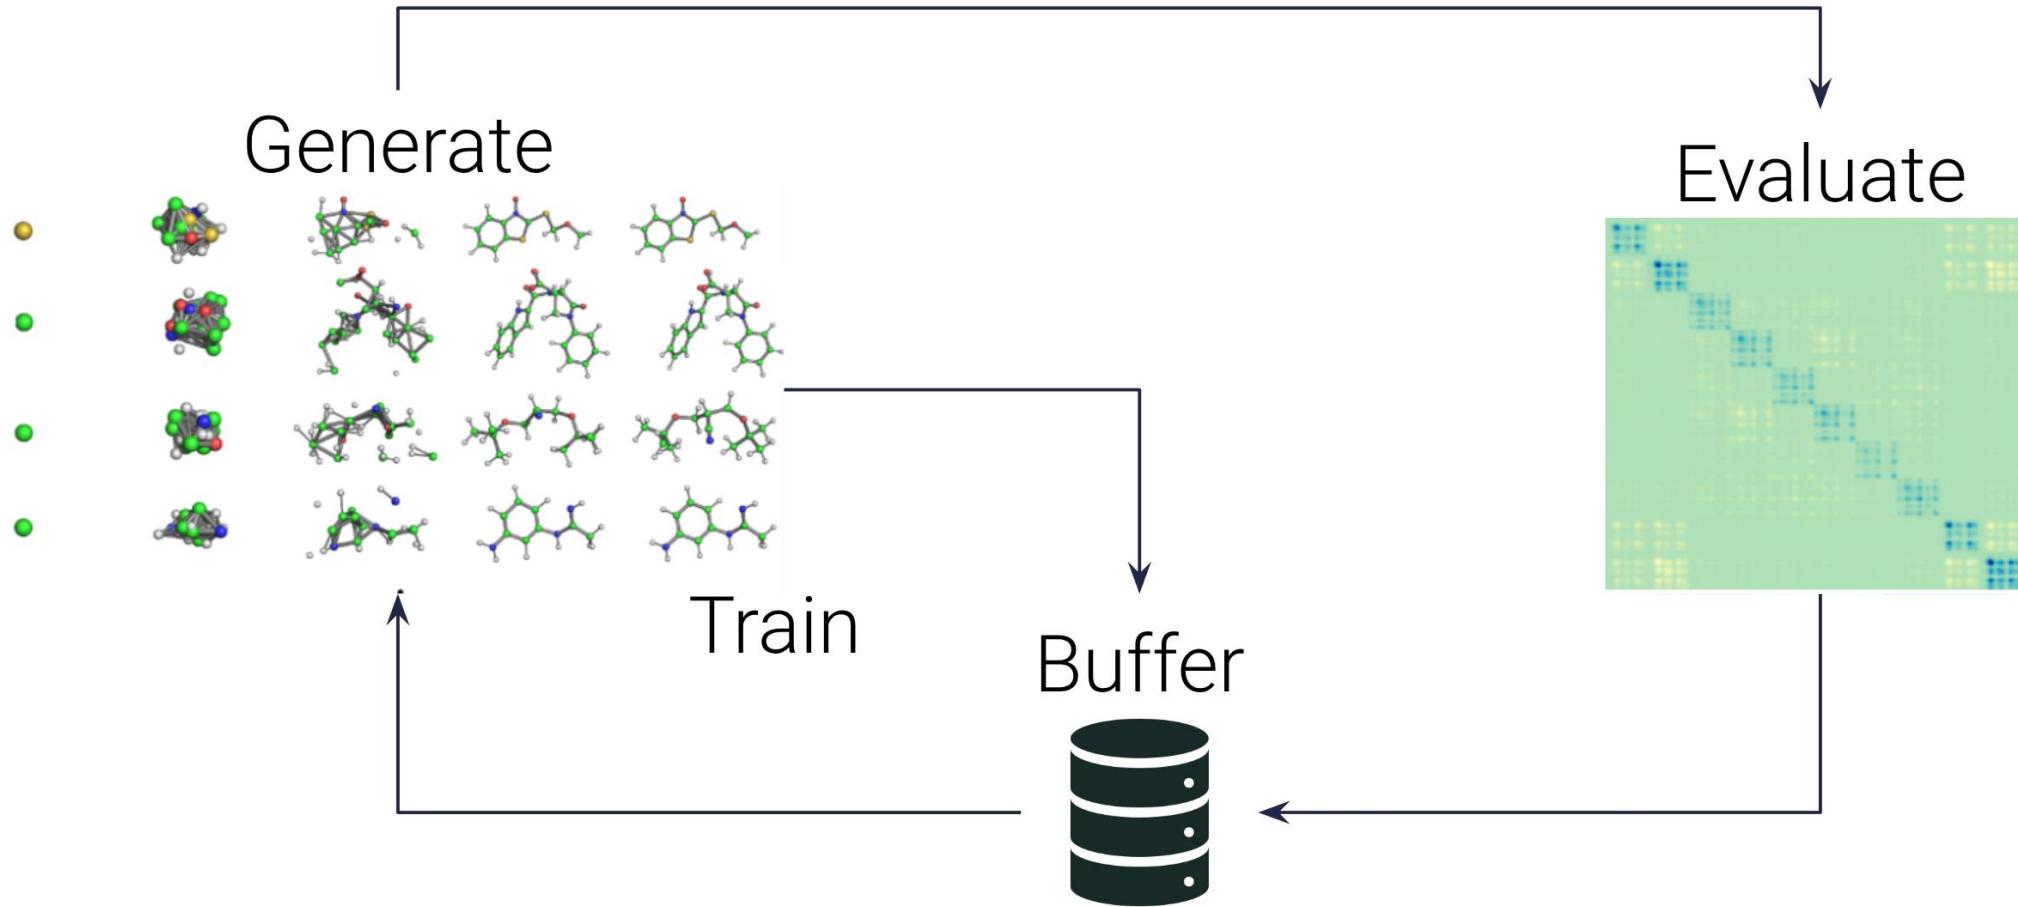
$$H_f(x, y) = \begin{pmatrix} \frac{\partial^2 f}{\partial x^2} & \frac{\partial^2 f}{\partial x \partial y} \\ \frac{\partial^2 f}{\partial y \partial x} & \frac{\partial^2 f}{\partial y^2} \end{pmatrix}$$

*Eigvals, eigvecs = torch.eigh(H_f)
Eigvals = [λ₀, λ₁, ...] # λ₀, λ₁ are the smallest two eigenvalues*

HIP predicts the Hessian from atom positions



We will train a generative model to generate transition states



We search for transition states using Gentlest Ascent Dynamics (GAD)

GAD equation:

$$\mathbf{x}' = -\nabla V(\mathbf{x}) + 2(\nabla V, \mathbf{v}_0)\mathbf{v}_0$$

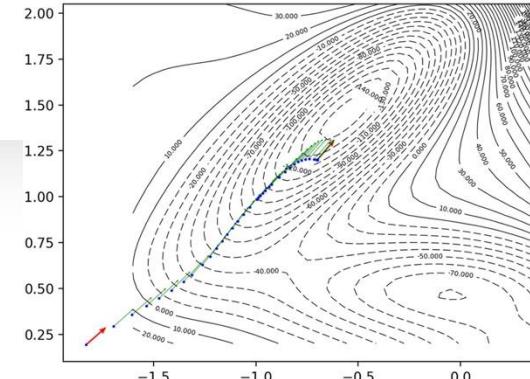
[1]

$V(\mathbf{x})$ = Energy scalar

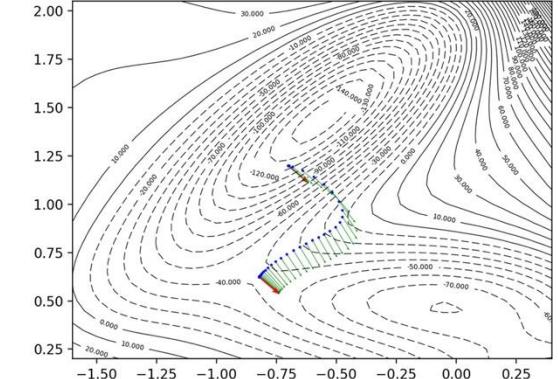
\mathbf{v}_0 = eigenvector of smallest eigenvalue

$\| \text{GAD step} \| = \| \text{Force} \|$

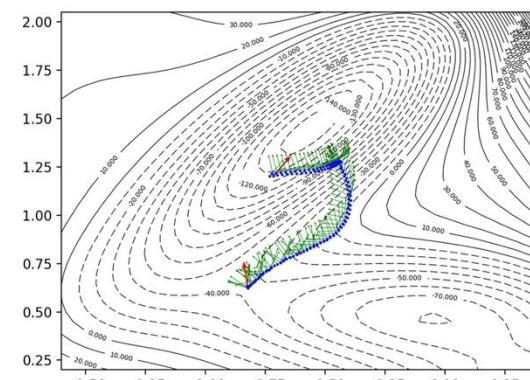
*First order
Integration
(Euler)



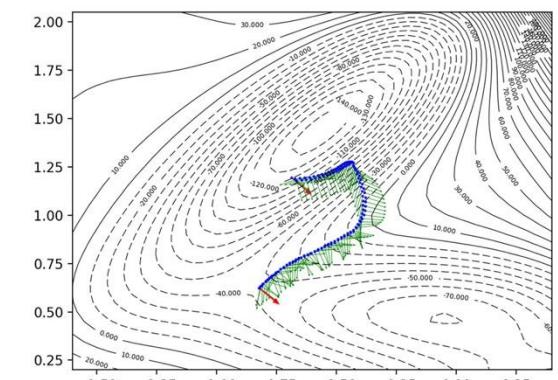
(a) GAD with \mathbf{v}_1 as initial control vector.



(b) GAD with \mathbf{v}_2 as initial control vector.



(c) GAD-CD with \mathbf{v}_1 as initial control vector.



(d) GAD-CD with \mathbf{v}_2 as initial control vector.

What GAD trajectories roughly look like

[1] W. E and X. Zhou, "The gentlest ascent dynamics," *Nonlinearity*, vol. 24, no. 6, pp. 1831–1842, May 2011, doi: <https://doi.org/10.1088/0951-7715/24/6/008>.

91.3% of trajectories starting from reactants found transition states

Experiment A:

GAD eq'n with Euler Integration (GAD-Euler)

Ending Geometries	Percentage
Minimum	2.8%
Transition State	91.3%
Unstable	3.8%
Higher Order Unstable	2.1%

$dt = 0.005\text{\AA}$, $n_samples = 300$, $max_steps = 500$

Average steps to convergence & average time per sample (RTX 2070):

<<Euler: **25.1 steps** to TS, **50s per sample**>>

<<RK45: **7.3 steps** to TS, **64s per sample**>>

GAD-Euler

Typical graph that ended at Transition State (91.3% of samples)

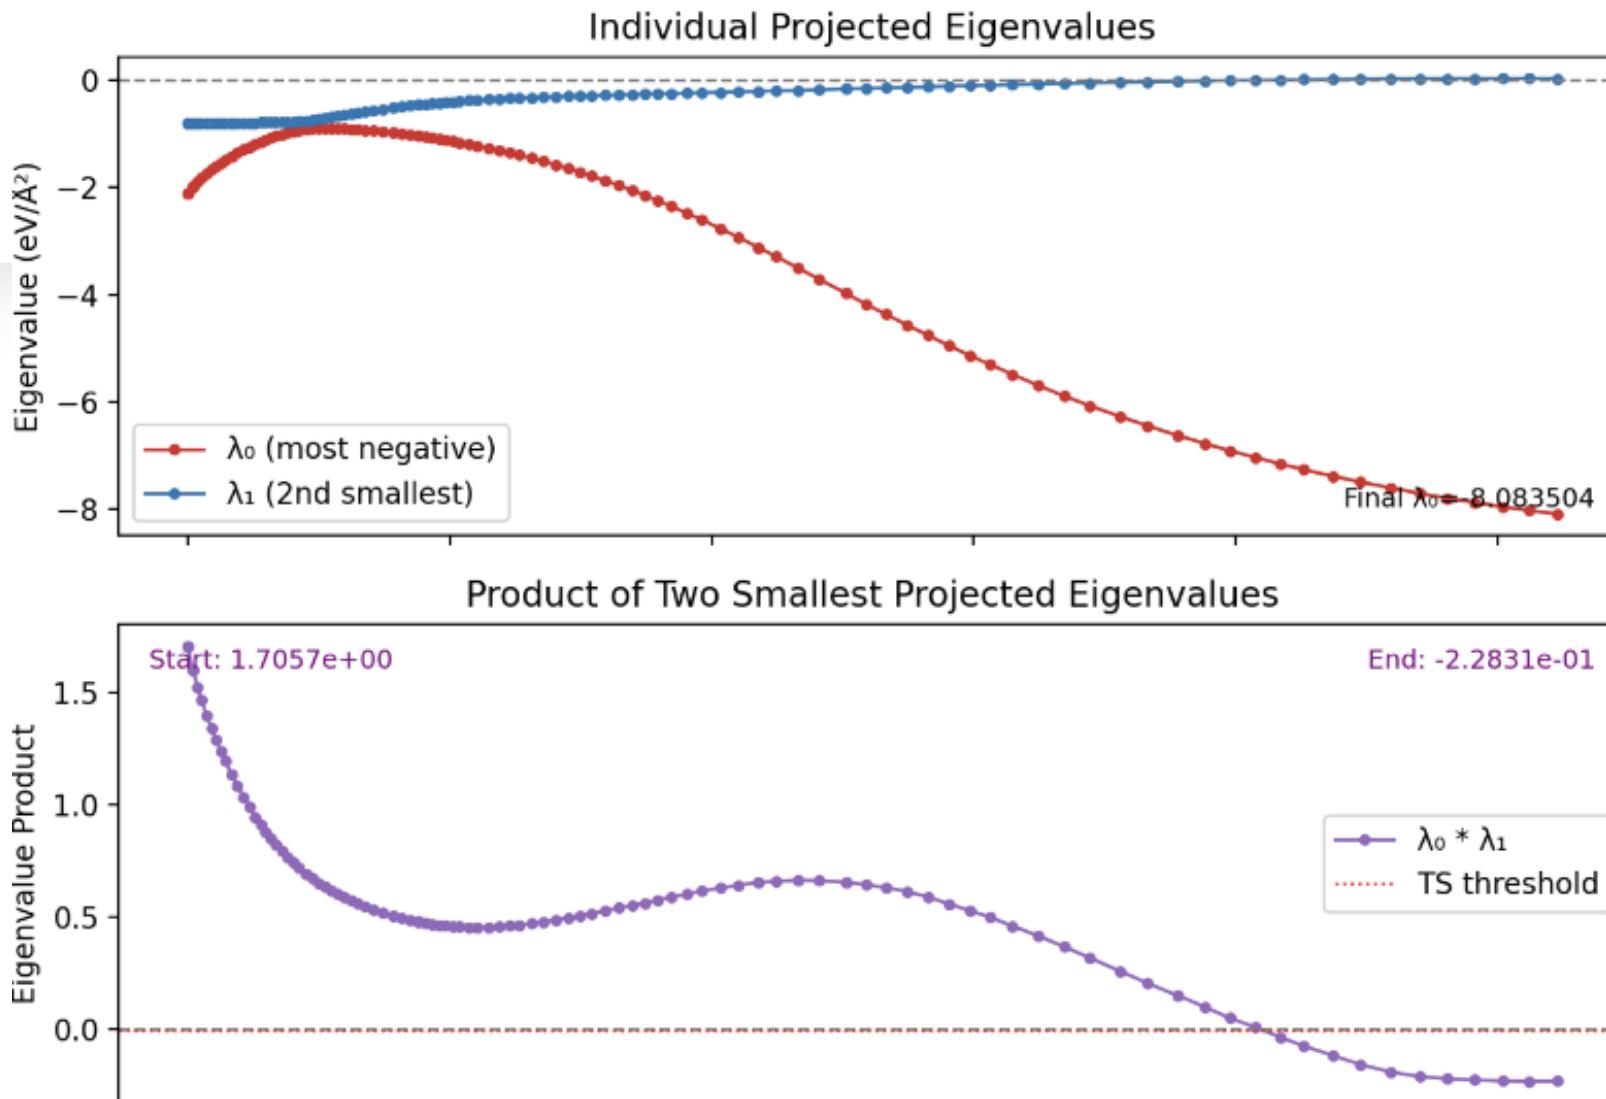


Figure 2: Trajectory of $\text{C}_2\text{H}_3\text{N}_3\text{O}_2$ starting, slowly converging to a TS

The samples that didn't reach TSs had two failure modes:

Stuck 1: 5.9% end up unstable

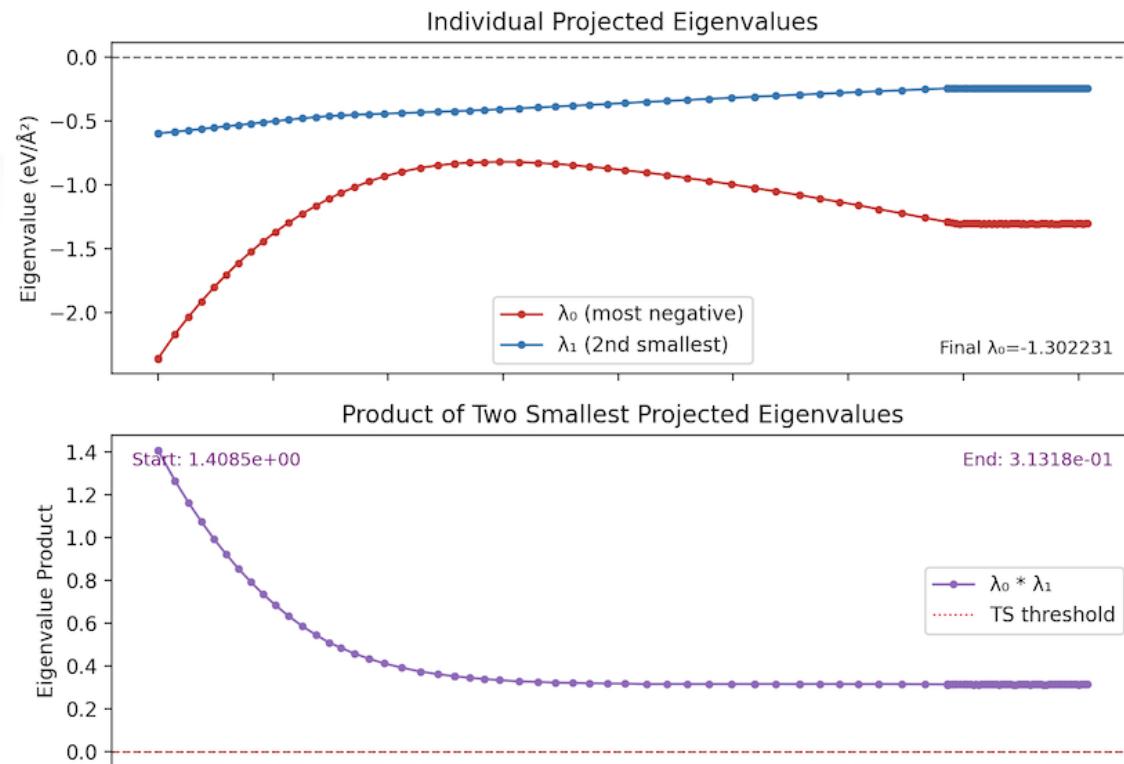


Figure 3: Trajectory of $\text{C}_2\text{H}_3\text{N}_3\text{O}_2$, unable to escape an unstable geometry

Stuck 2: 2.8% end up unstable

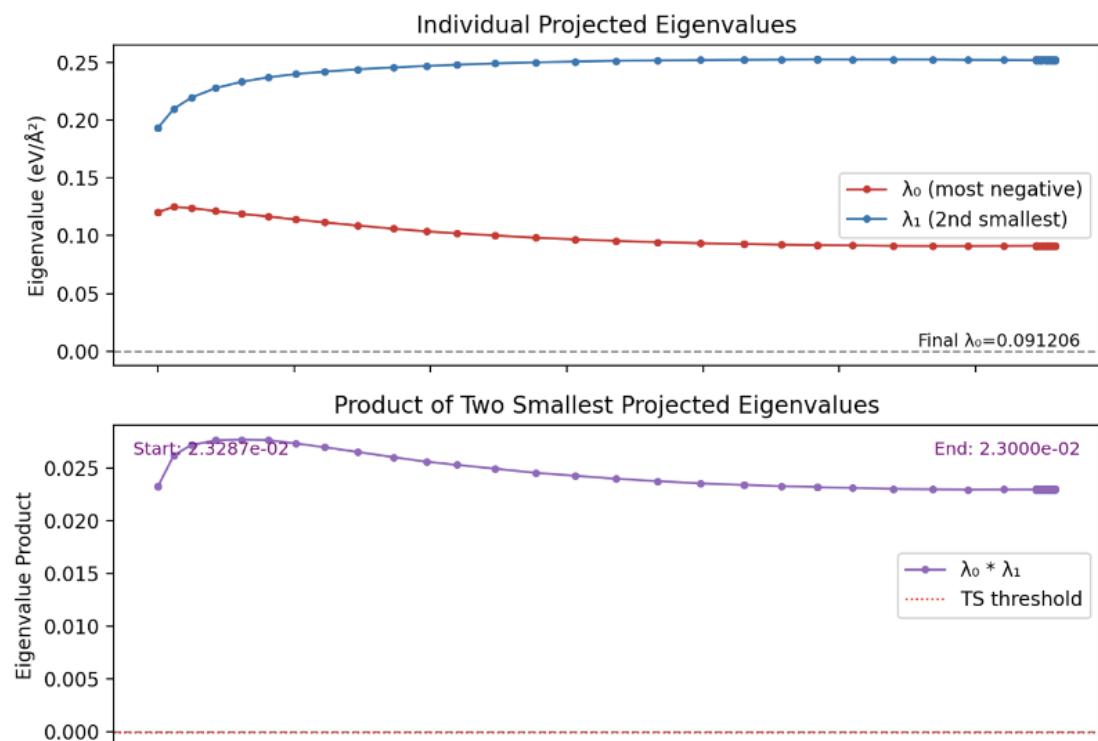


Figure 3: Trajectory of $\text{C}_2\text{H}_3\text{N}_3\text{O}_2$, unable to escape an energy minimum

We've confirmed GAD works.
Can we try to directly optimize λ_0 and λ_1 ?

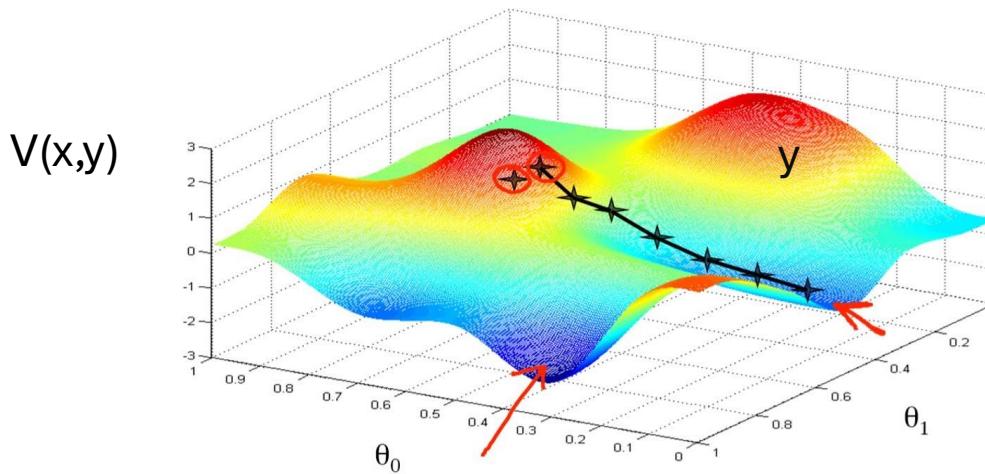
Optimization

Core idea:

Directly find gradients of λ_0 , λ_1 , and just
optimize towards a Transition State ($\lambda_0^* \lambda_1$).

Benefits:

- Much simpler algorithm
- Could lead to more direct paths than GAD
- Could be faster than GAD



The most direct way of approaching our goal (TS) is to run a gradient descent on the TS condition directly

Optimization

TS Condition:
negative $\lambda_0^* \lambda_1$



Gradient descent on:
 $\lambda_0^* \lambda_1$ until it converges to
negative

Gradient path: loss → eigenvalues → projected Hessian → raw Hessian → HIP model → coordinates

The most direct way of approaching our goal (TS) is to run a gradient descent on the TS condition directly

Optimization

```
for step in range(n_steps):
    # Compute Hessian and eigenvalues
    hessian = model.get_hessian(coords)
    hessian_proj = project_out_rotations_translations(hessian, coords)
    eigenvalues = torch.linalg.eigvalsh(hessian_proj)

    # Loss: product of two smallest eigenvalues
    λ0, λ1 = eigenvalues[0], eigenvalues[1]
    loss = λ0 * λ1 # Negative at TS (λ0 < 0, λ1 > 0)

    # Gradient descent
    grad = torch.autograd.grad(loss, coords)[0]
    coords = coords - lr * grad

    # Early stop if product is sufficiently negative
    if loss < -1e-4:
        break
```

92.7% of samples converged (starting from reactant position)

Eigenvalue Product Descent

Experiment B:

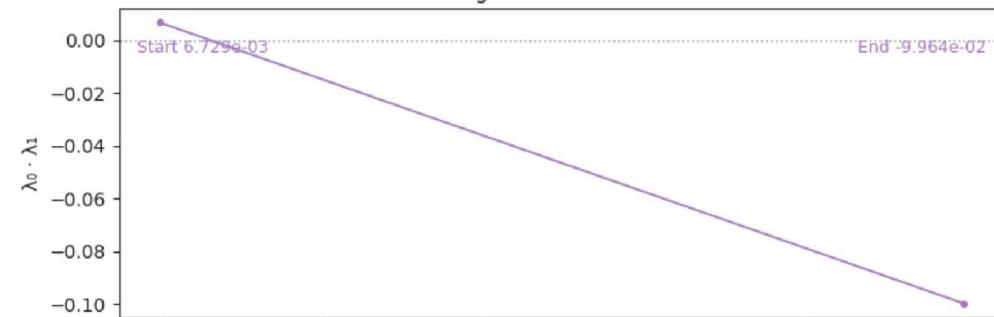
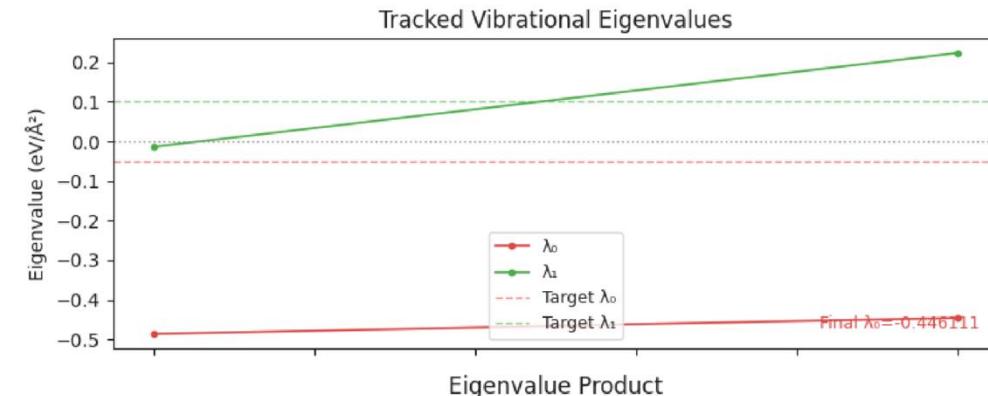
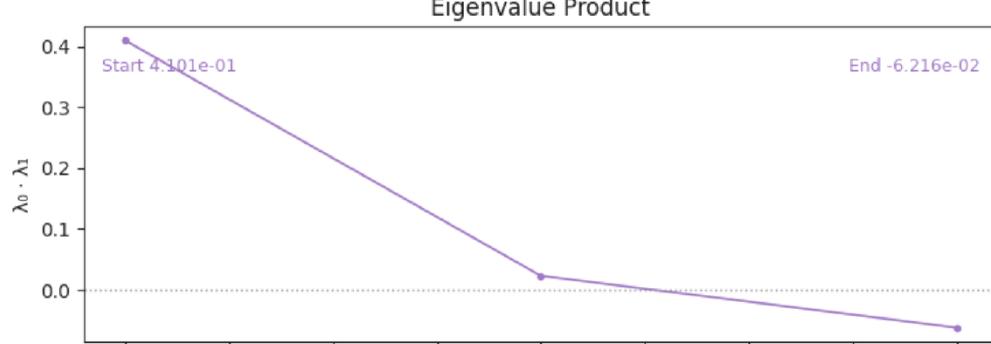
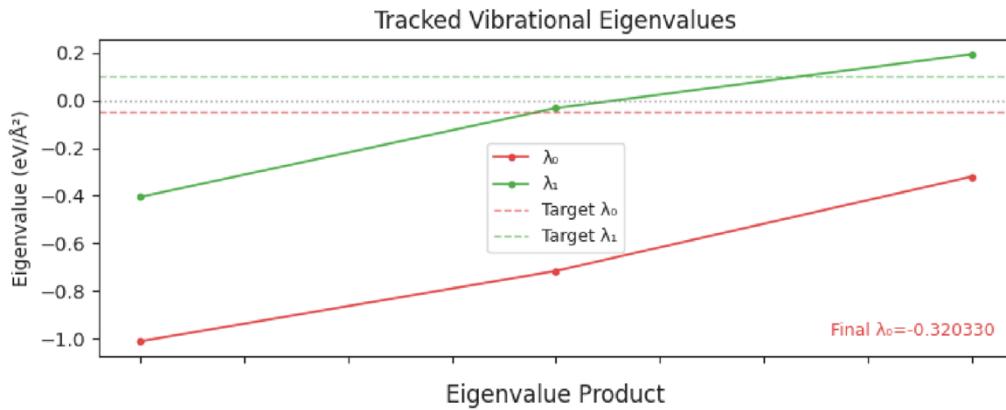
Eigenvalue product descent with Euler stepping

Loss Function	TS Success Rate (%)	Avg λ_0 (eV/Å²)	Avg λ_1 (eV/Å²)	Avg $\lambda_0 \times \lambda_1$
Product Descent	92.7	-0.2405	0.0763	-8.260e-03

$dt = 0.01 \text{ Å}$, $n_samples = 300$, $max_steps = 500$

The samples converged almost immediately (within 1-3 steps)

Eigenvalue Product
Descent



Average steps to convergence & average time per sample (RTX 2070):

<<Euler: **25.1 steps** to TS, **50s per sample**>>

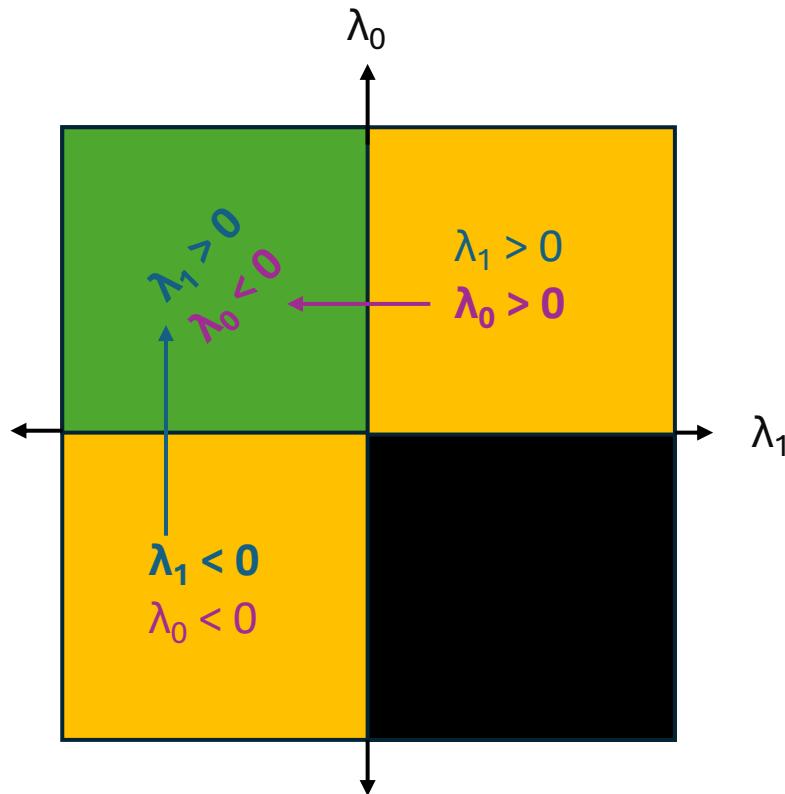
<<RK45: **7.3 steps** to TS, **64s per sample** >>

<<**EigProd: 2.3 steps to TS, 3.2s per sample** >>

What if there is an EVEN MORE direct way of approaching transition states

Direct Descent

*Only change λ_0 OR λ_1 , one at a time,
depending on the quadrant you are in*



What if there is an EVEN MORE direct way of approaching transition states

Direct Descent

```
for step in range(n_steps):
    # Compute Hessian and eigenvalues
    hessian = model.get_hessian(coords)
    hessian_proj = project_out_rotations_translations(hessian, coords)
    eigenvalues = torch.linalg.eigvalsh(hessian_proj)

    n_negative = (eigenvalues < 0).sum()
    λ0 = eigenvalues[0]

    # Adaptive loss based on current saddle order
    if n_negative == 0:
        # No negatives → push λ0 below target
        loss = (λ0 - neg_target) ** 2
    elif n_negative == 1:
        # Exactly one negative → done!
        loss = 0.0
        break
    else:
        # Too many negatives → push extras positive
        loss = ((pos_floor - eigenvalues[1:]) ** 2).sum()

    # Gradient descent
    grad = torch.autograd.grad(loss, coords)[0]
    coords = coords - lr * grad
```

Direct Descent Performed better,
and was slightly faster

Direct Descent

Loss Function Comparison (n=287, 500 steps)

TS Success = Converged to exactly 1 negative eigenvalue

Loss Function	TS Success Rate (%)	Avg λ_0 (eV/Å²)	Avg λ_1 (eV/Å²)	Avg $\lambda_0 \times \lambda_1$
Direct Descent	99.7	-0.1528	0.1241	-1.733e-02
Product Descent	92.7	-0.2405	0.0763	-8.260e-03

Average steps to convergence & average time per sample (RTX 2070):

<<Euler: **25.1 steps** to TS, **50s per sample**>>

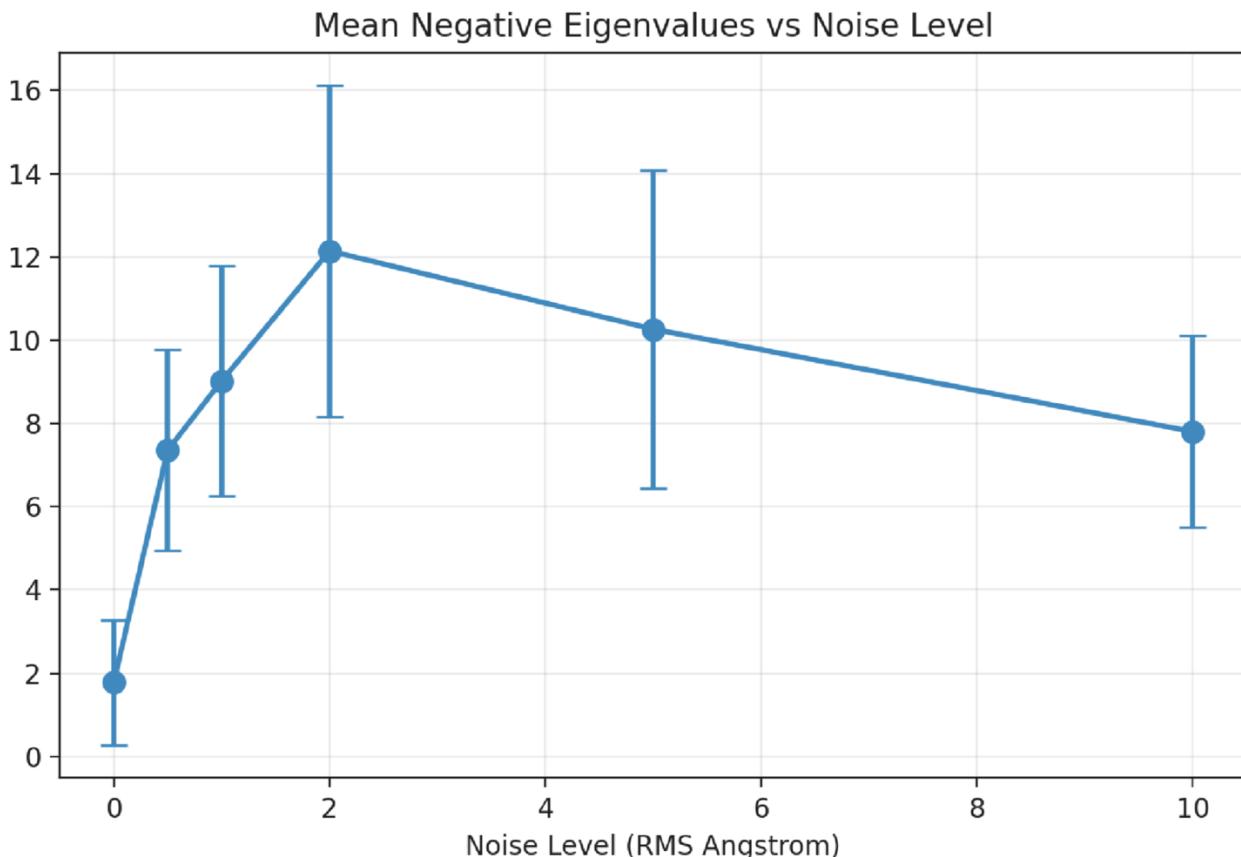
<<RK45: **7.3 steps** to TS, **64s per sample** >>

<<EigProd: **2.3 steps** to TS, **3.2s per sample** >>

<<Direct: **1.8 steps** to TS, **3.1s per sample** >>

**So far, we've started from reactants.
But the generative model will start from random
geometries.**

Noisy Experiments



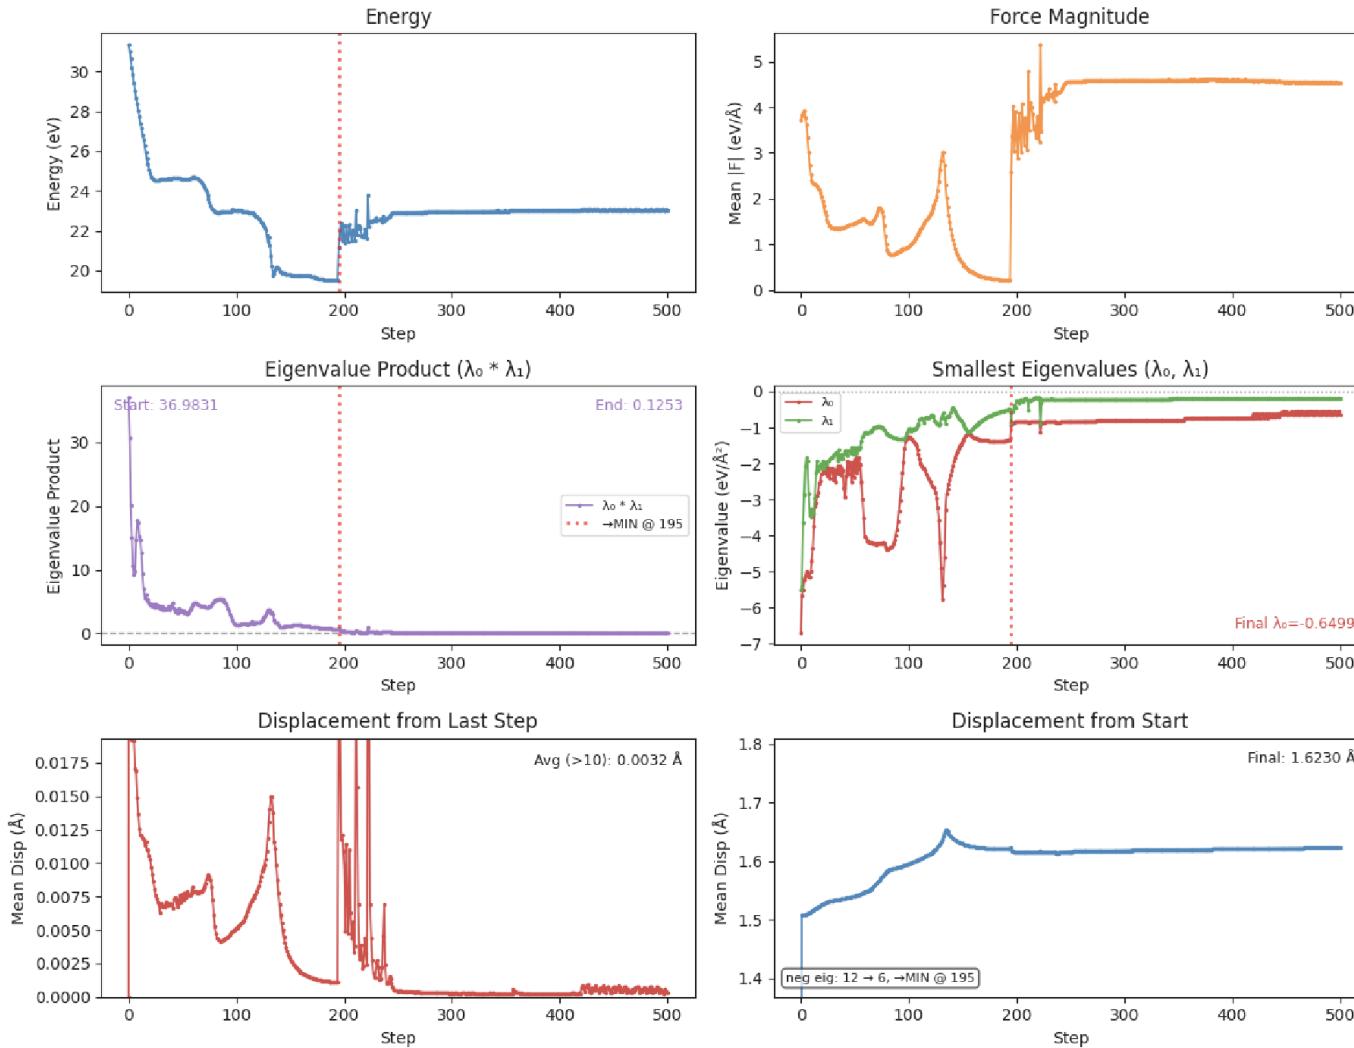
- We want random starts with higher # of imaginary frequencies, as that is exactly what the generative model will produce.

How many negative eigenvalues we get based on how much noise we add

Optimization

Pure GAD almost always plateaued at order 4-8 negative eigenvalues (13% convergence rate to TS)

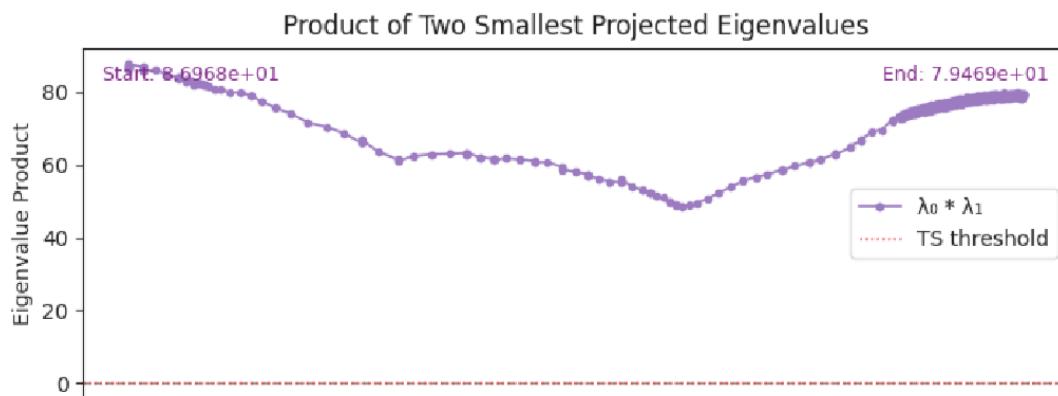
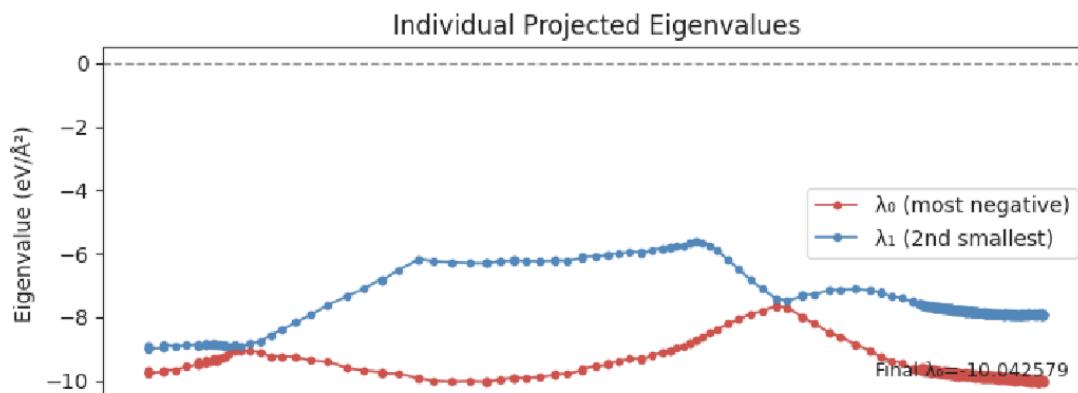
Noisy Experiments –
GAD-Euler



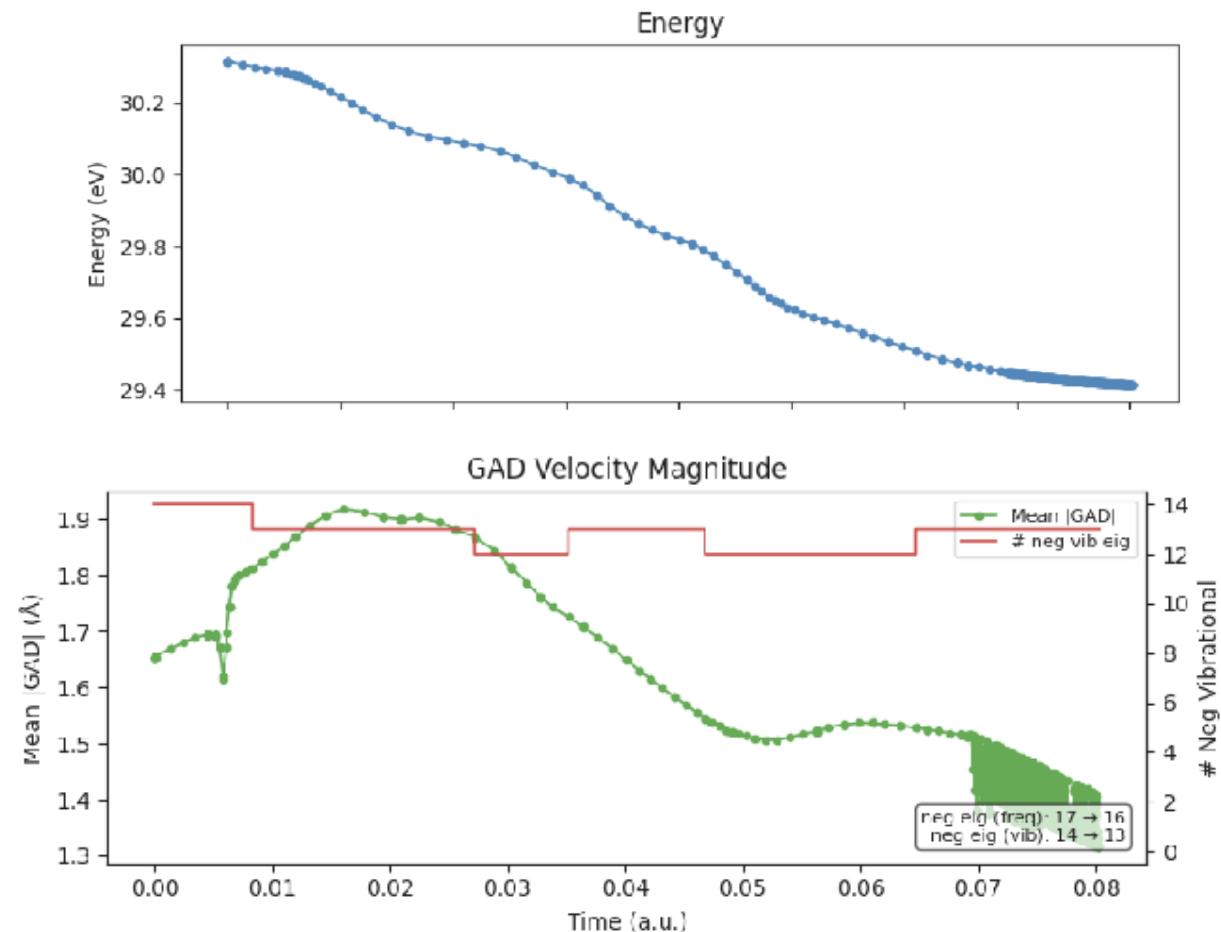
Typical GAD trajectory,
which gets stuck

GAD-RK45 also plateaued, but sometimes oscillated (example 1) (<5% convergence rate to TS)

17 negative eigenvalues to 16



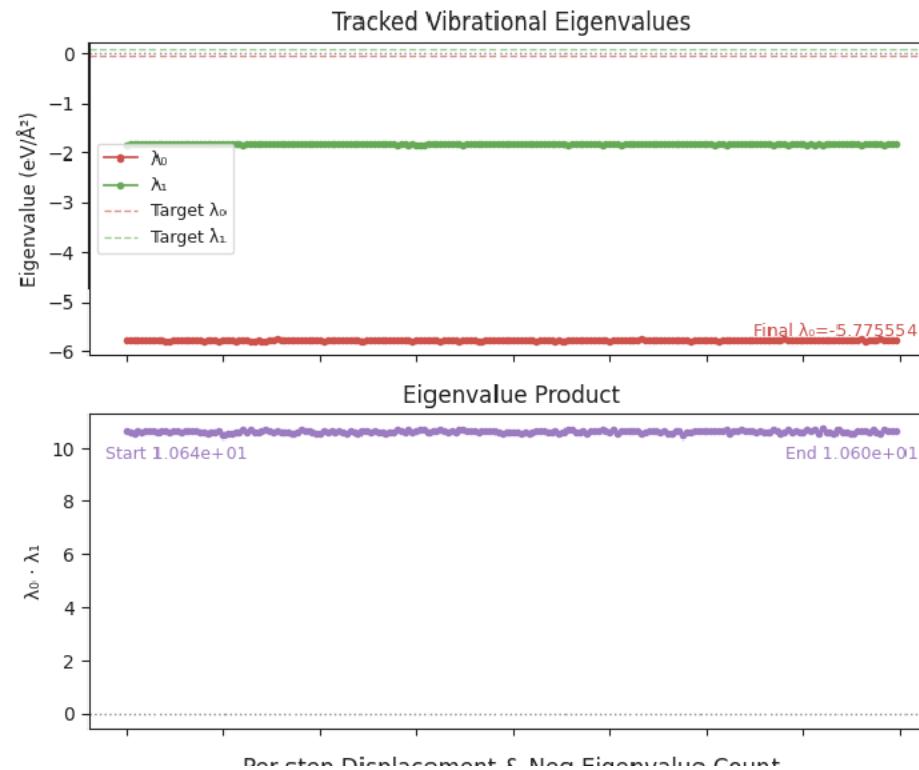
Noisy Experiments – GAD RK45



Direct Descent Methods had almost no effect on stability, and none converged

Noisy Experiments –
Direct Descent

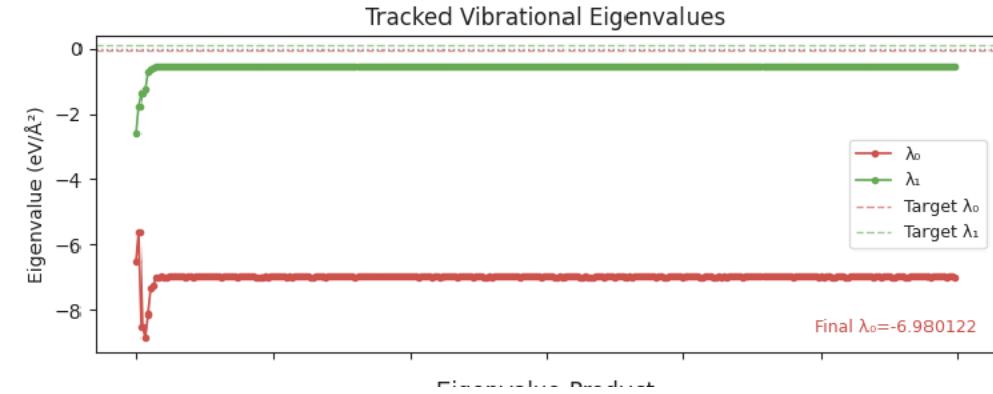
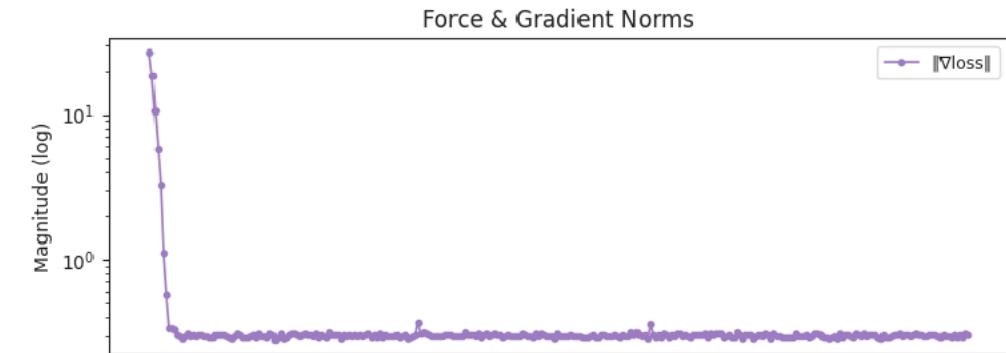
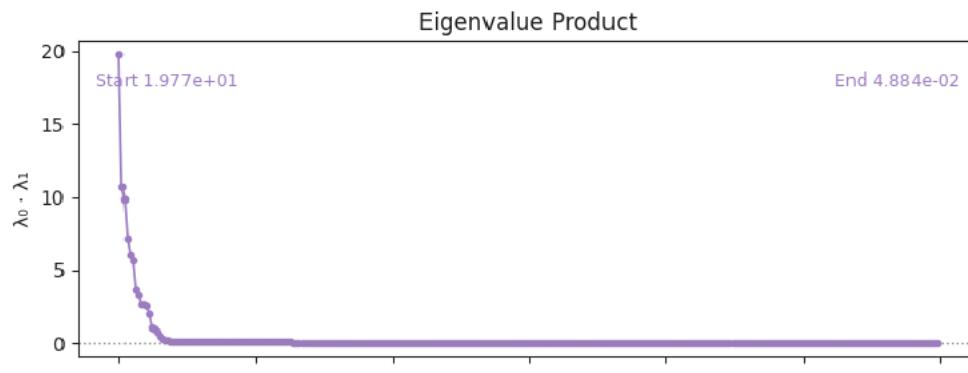
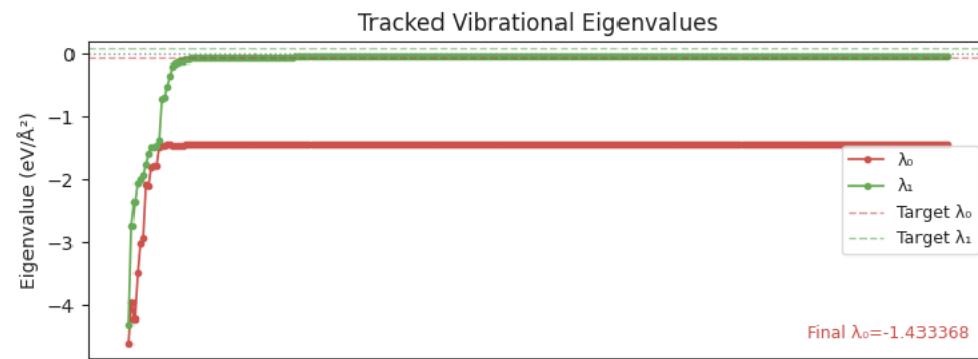
Nearly all graphs resemble the following:



**Product Descent very quickly reached plateaus,
<5% converged**

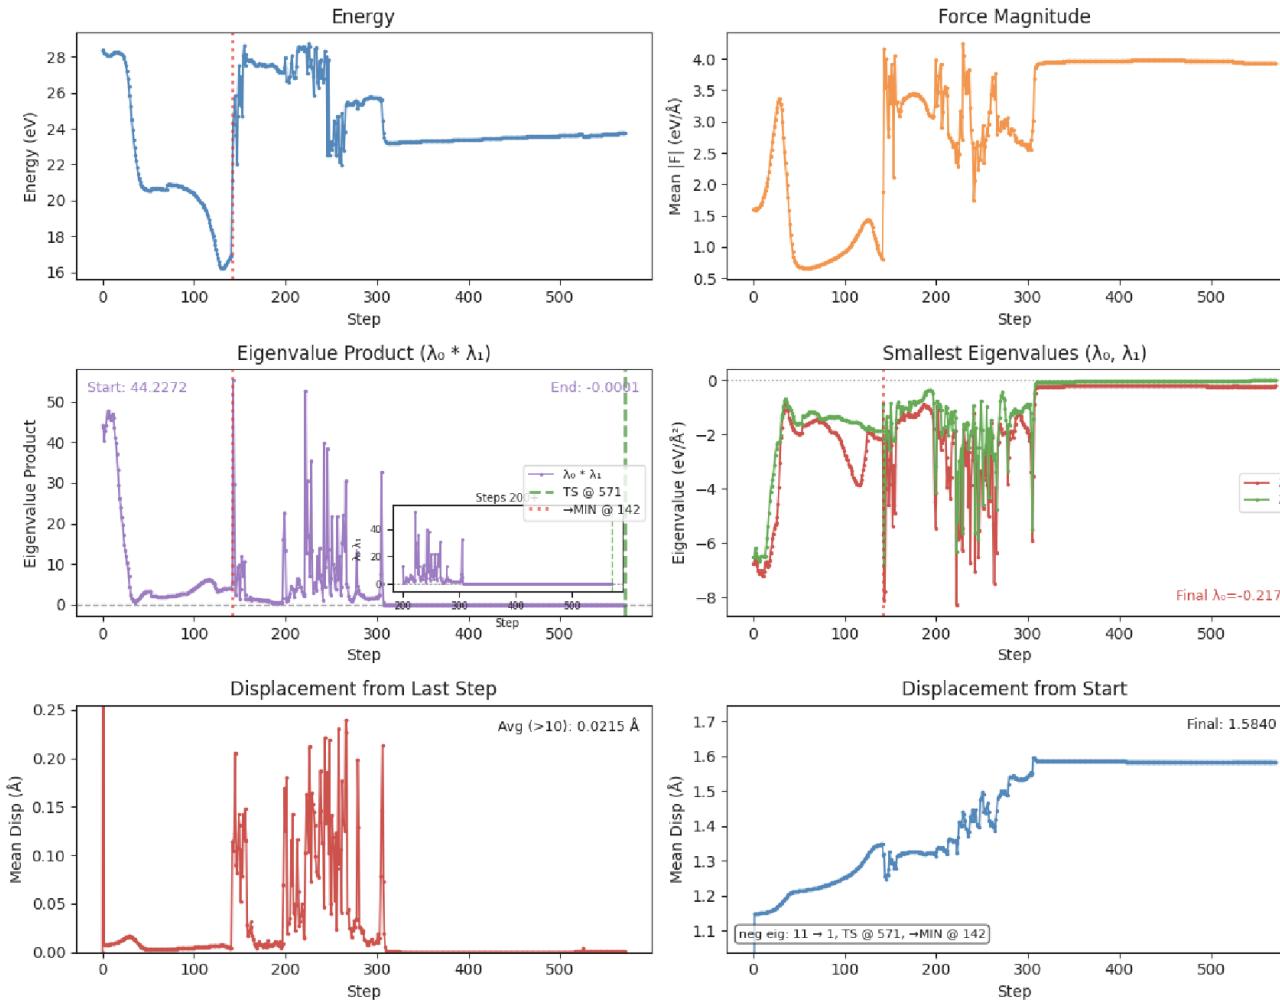
Noisy Experiments –
Product Descent

Nearly all graphs resemble the following:



While many hybrid methods didn't work, GAD + Product Descent improved performance! (~33% Convergence) (Converged Example)

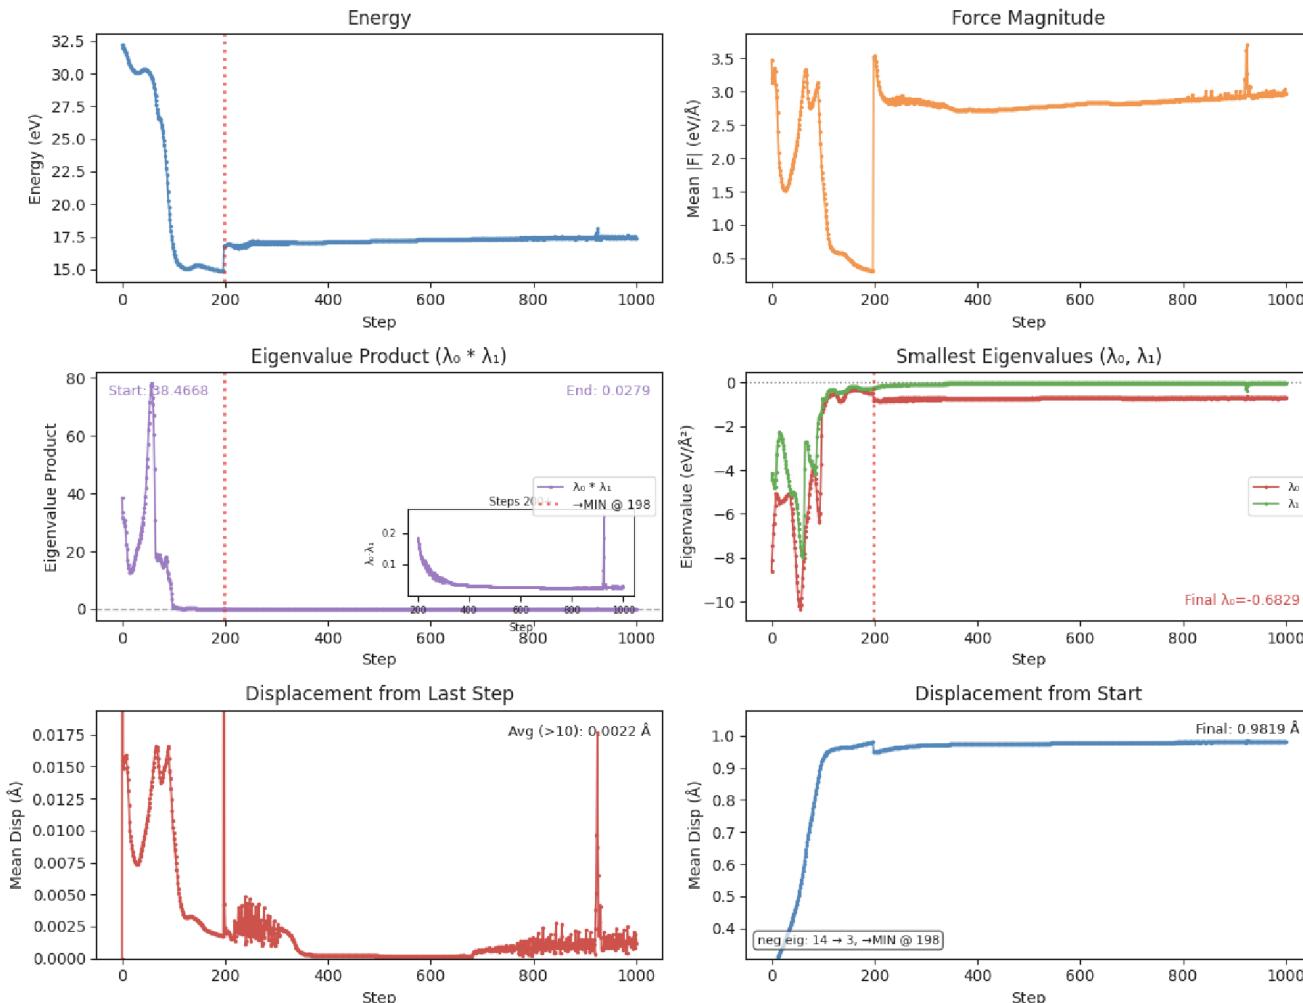
Noisy Experiments –
Hybrid Method



```
# --eig-descent-lr 0.01
# --eig-descent-max-step 0.5
# --stop-at-ts
# --n-steps 1000
# --dt 0.005
```

While many hybrid methods didn't work, GAD + Product Descent improved performance! (~33% Convergence) (Stalled example)

Noisy Experiments – Hybrid Method



```
# --eig-descent-lr 0.01
# --eig-descent-max-step 0.5
# --stop-at-ts
# --n-steps 1000
# --dt 0.005
```

Quick note about GAD vs. Eigenvalue Descent

- GAD Cannot Escape Reactant Geometries (Direct Descent is still king).

Noisy Experiments –
Hybrid Method

[Transition Distribution]
0neg-to-0neg: 5 samples
1neg-to-1neg: 28 samples
2neg-to-1neg: 49 samples
3neg-to-1neg: 18 samples

New analytical forcefield: SCINE DFTB0

GAD Experiments

- Doesn't preserve gradients (

```
[Transition Distribution]
0neg-to-0neg: 5 samples
1neg-to-1neg: 28 samples
2neg-to-1neg: 49 samples
3neg-to-1neg: 18 samples
```

HIP Enables Efficient Transition State Discovery

- Confirmed that GAD is a viable tool for TS search.
- Direct eigenvalue descent with access to gradients is a better approach
- (expected, as it's literally optimizing our criteria, and the paths are not extremely complicated.)
- Can only achieve ~33% convergence on noisy geometries

December Experiments

GAD Algorithm Improvements for Transition State Search

From Noisy Geometries to Reliable Convergence

Overview

1. **New Forcefield Calculator** (SCINE DFB0)
2. **L-BFGS Energy Minimizer** — Reduce saddle order first
3. **Plateau Detection (Kicking)** — Adaptive step sizes
4. **Higher-Order GAD** — Multi-mode escape
5. **Mode Tracking + Trust Radius** — Final solution (92% / 71%)

SCINE Sparrow is a different level of theory than T1x and HIP

Key Insight

HIP and SCINE forcefields disagree significantly on the number of negative eigenvalues, especially for reactant geometries.

Table 1: Summary of Mean Negative Eigenvalues by Geometry and Calculator
($n = 100$)

Geometry	HIP Mean	SCINE Mean	Agreement
midpoint_rt	3.55	4.84	23 (23.0%)
reactant	0.00	1.80	5 (5.0%)
product	0.10	2.18	10 (10.0%)
midpoint_rt_noise1A	9.55	11.36	11 (11.0%)
midpoint_rt_noise2A	13.42	10.07	10 (10.0%)

Key Observations:

- **95% disagreement** on reactant geometries
- SCINE only sees 5% of reactants that HIP sees
- **Fix:** Apply climbing algorithm when minima detected

**(In any other geometry,
HIP and SCINE are
quite consistent.)**



<https://github.com/qcscine/sparrow>

Like HIP, starting from reactants, SCINE can achieve 100% convergence using GAD with some simple tweaks

Key Insight

Pure GAD achieves **100% convergence on non-minima** samples. The challenge is making it work on noisy geometries with many negative eigenvalues.

Overall Statistics (Pure GAD on TS geometries)

λ_0 : $-0.074139 \pm 0.079542 \text{ eV}/\text{\AA}^2$
 λ_1 : $0.044880 \pm 0.062035 \text{ eV}/\text{\AA}^2$
 $\lambda_0 \cdot \lambda_1$: $-5.443389e-03 \pm 7.901079e-03$
Steps taken: 18.2 ± 28.8
Steps to TS: 18.2 ± 28.8
Final time: $0.132 \pm 0.182 \text{ s}$

Transition Distribution

Oneg-to-1neg: 5 samples
1neg-to-1neg: 11 samples
2neg-to-1neg: 14 samples

While these didn't converge automatically with GAD, a simple kick was applied to get them out of minima.



Noisy Geometry Experiments

Where we were:

- Previous attempts either go in wrong directions, or plateau
- Cannot escape highly unstable regions

3 – Plateau Detection (Kicking)

Instead of directly using our algorithms, what if we just minimize the energy (*and possibly fine-tune later using our algorithms*)?

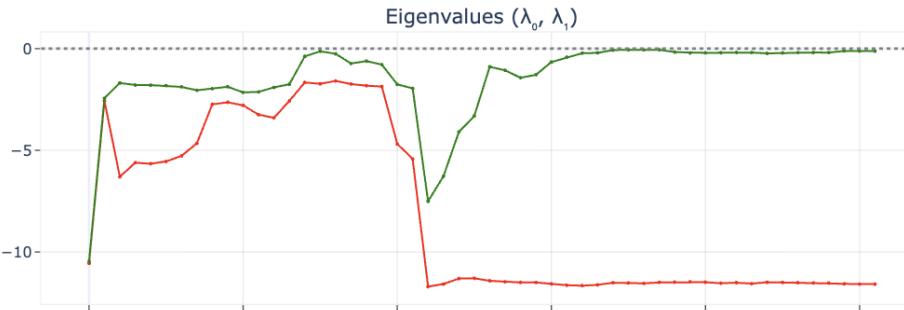
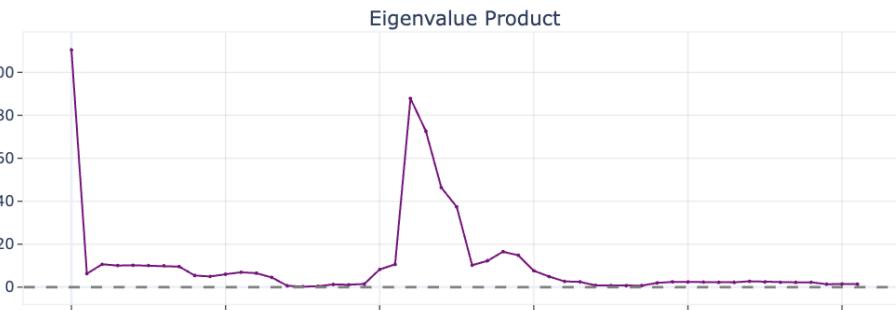
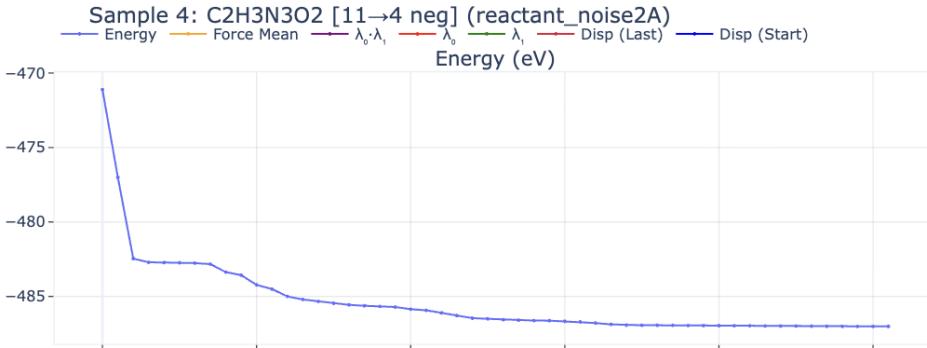
Key Insight

Idea: Since GAD converges 100% from non-minima, first reduce the saddle order using L-BFGS energy minimization, then apply GAD.

The Algorithm

1. Project forces into vibrational subspace (mass-weighted Eckart projection)
2. Use projected forces as gradients for L-BFGS
3. Check eigenvalue count periodically
4. Stop when target saddle order reached (e.g., 1 negative eigenvalue)

While the order of instability is reduced in most samples, they plateau, and there are many who have very little effect on instability.



SCINE Results

10/100 Convergence

- Pretty good reduction of order
- 2.5s–10s per sample
- Good reduction if starting above 10 eigenvalues
- Many plateau at ~ 8 negative eigenvalues

HIP Results

0/100 Convergence

- Number of neg-eigs plateau
- Energy changes become negligible
- Really good reduction of order
- 0.8s–4.5s per sample

Conclusion: L-BFGS alone is insufficient. Need a more elegant approach.

4 – Plateau Detection (Kicking)

GAD is the most elegant way forward, but it plateaus (stepsize becomes tiny). What if we kick it in the GAD direction by a big step?

Key Insight

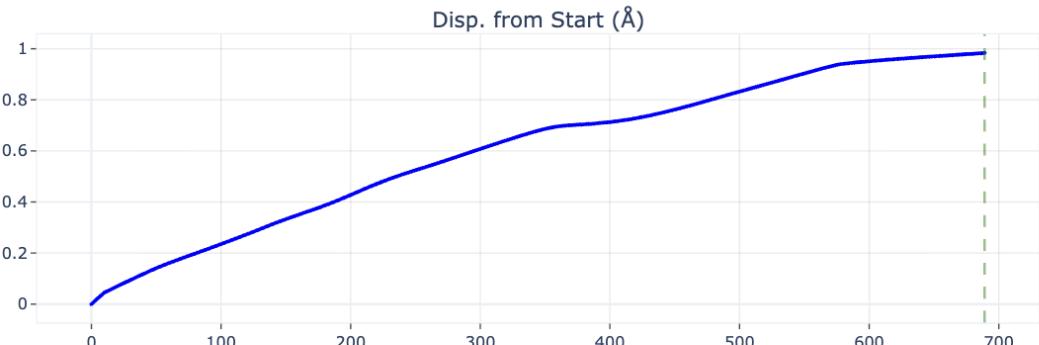
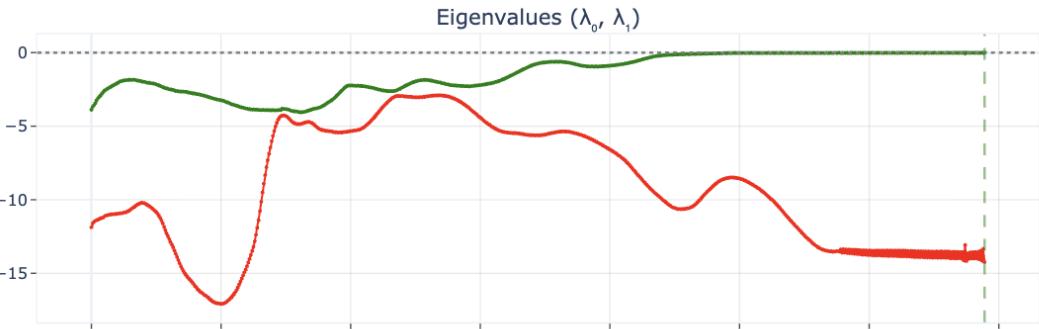
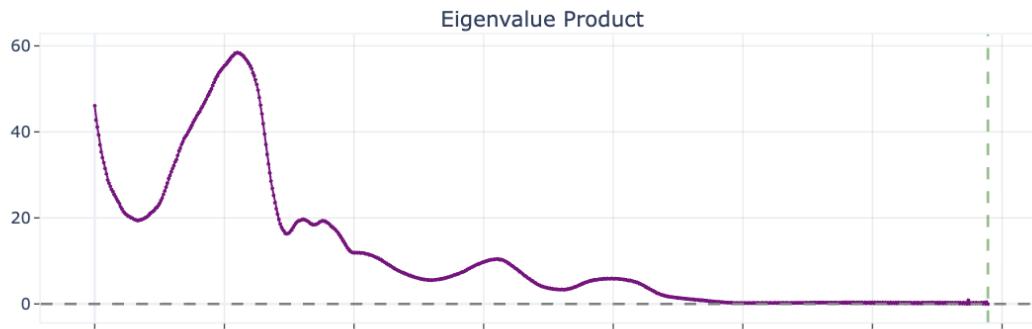
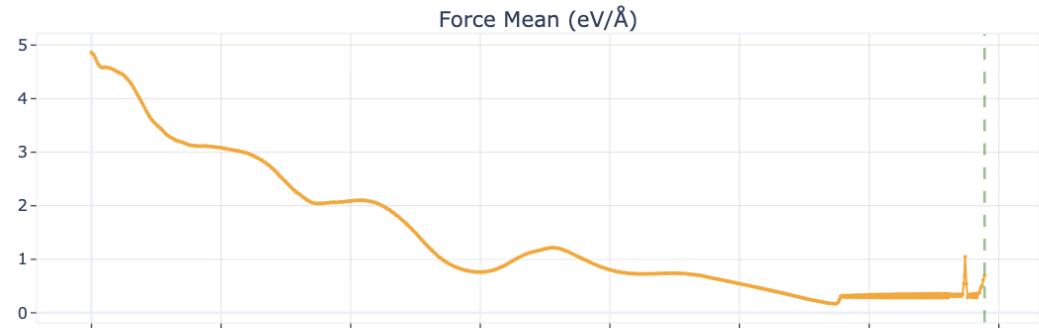
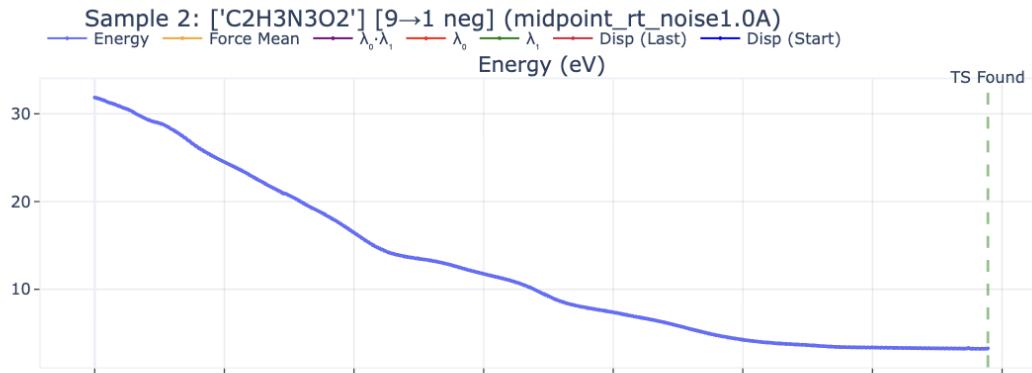
During GAD optimization, the algorithm can get “stuck” at high-index saddle points.

Solution: Monitor negative eigenvalue count and adaptively adjust step size.

4.1 The Algorithm

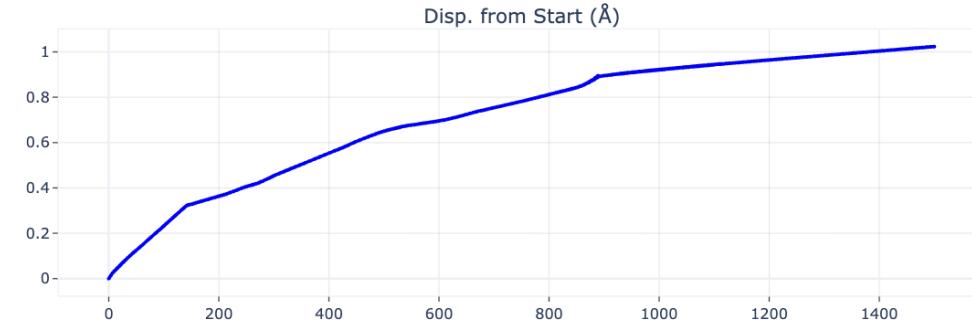
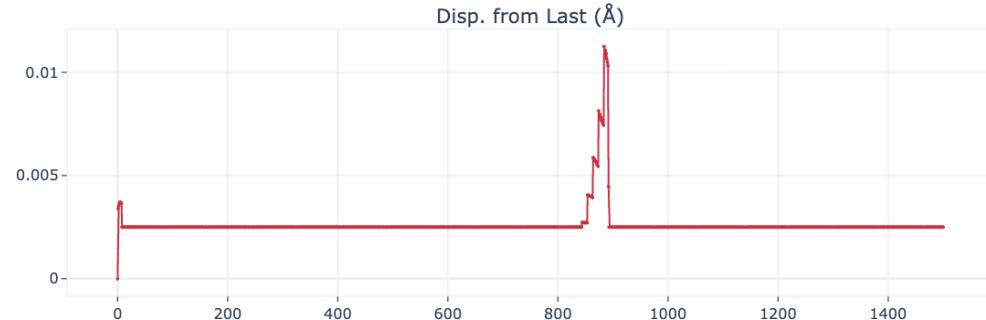
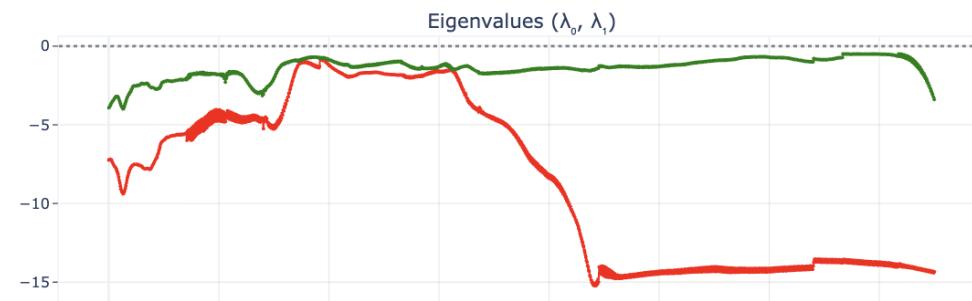
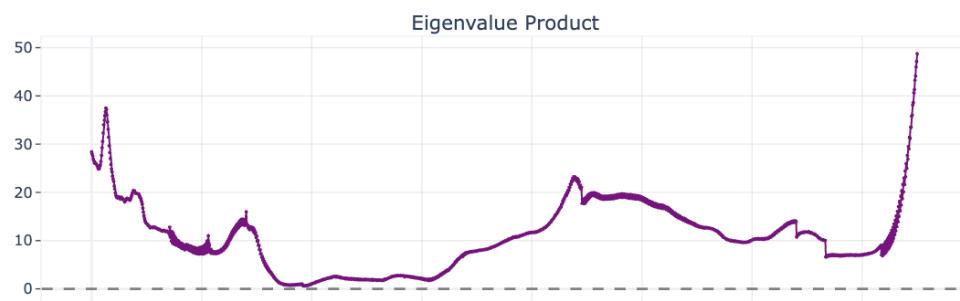
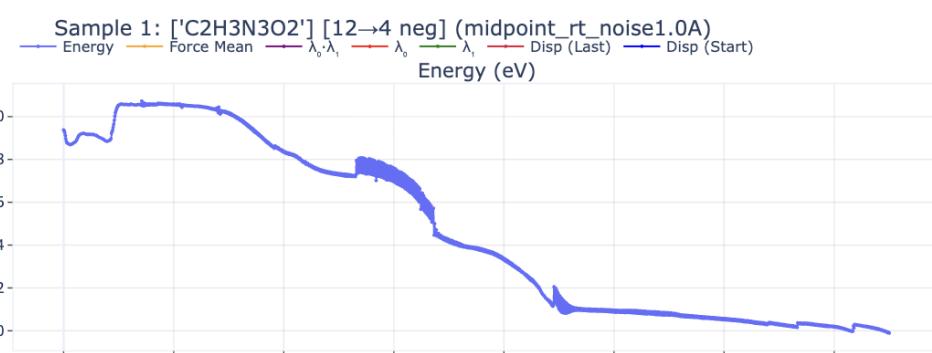
1. Track best_neg_vib (lowest saddle order seen)
2. **If improved:** Reset step size to base dt
3. **If worsened:** Shrink step size (more careful steps)
4. **If unchanged for patience steps:** Boost step size (kick!)
5. Enforce minimum step size floor ($500 \mu\text{\AA}$)

Successful sample that was saved by step size floor + kicking

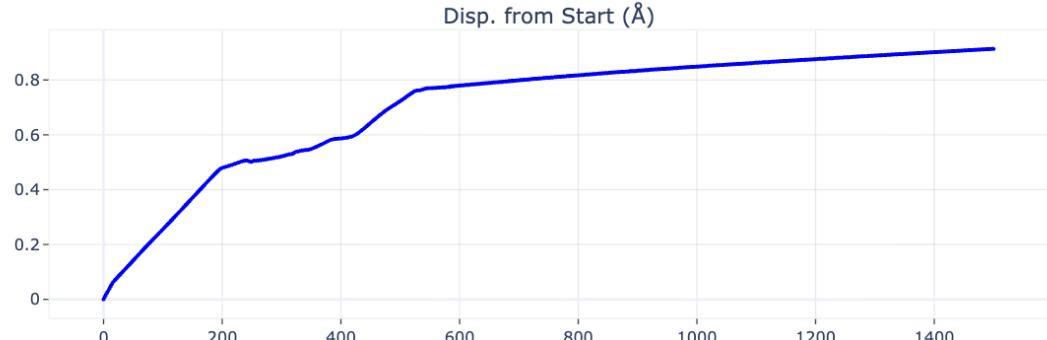
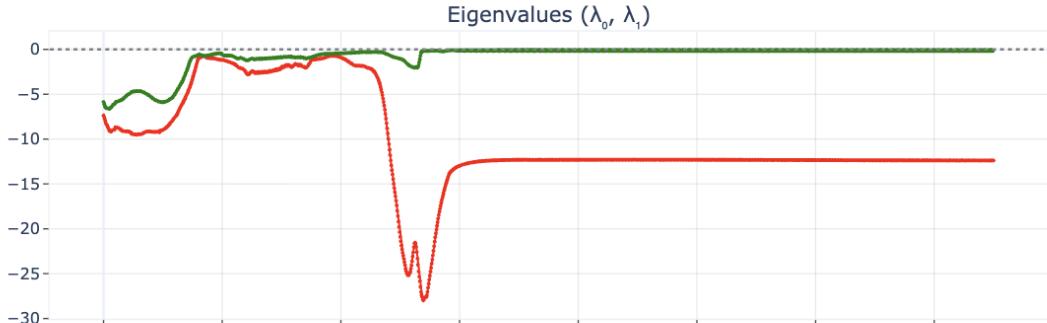
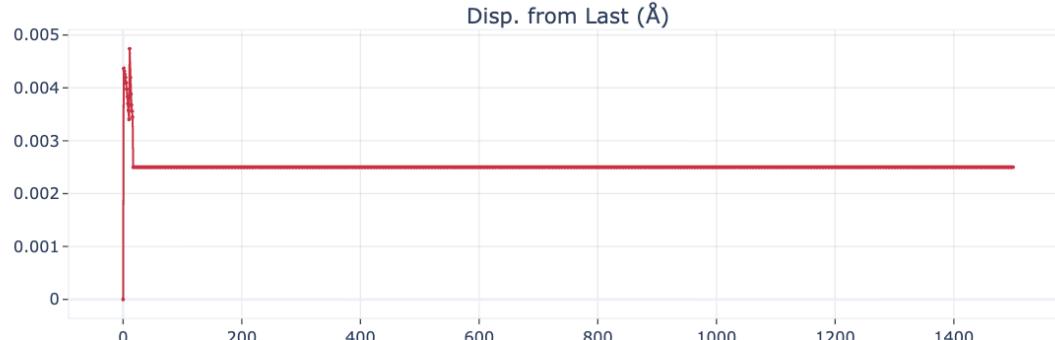
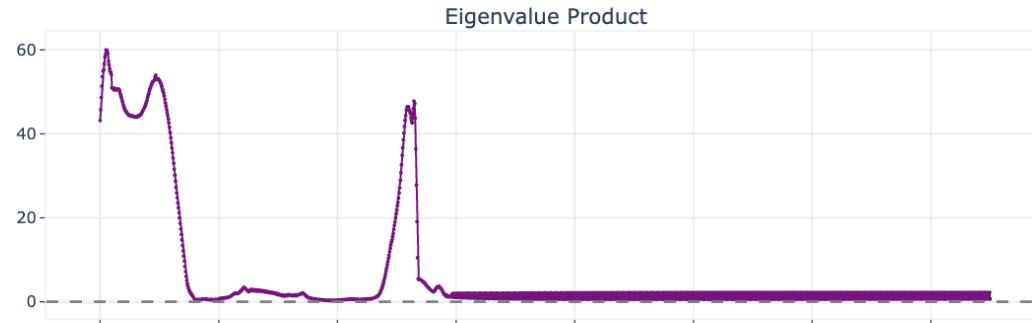
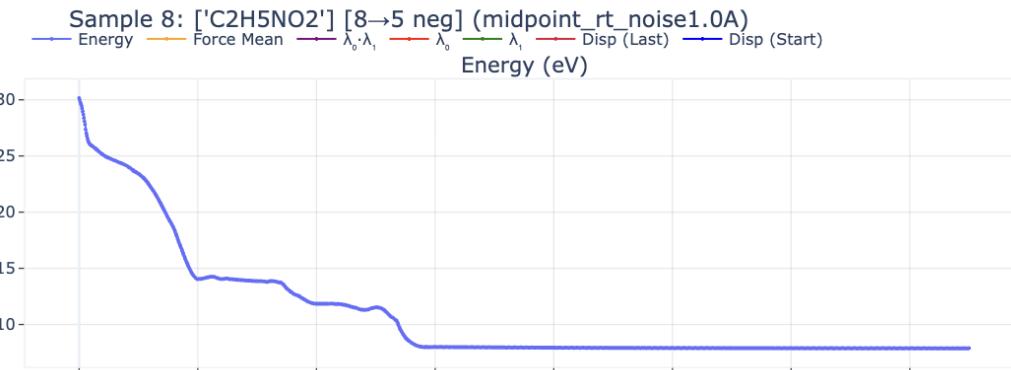


4 – Plateau Detection (Kicking)

Sample that diverged because of kicking



Sample that started oscillating as it stepped with size stepsize floor



4 – Plateau Detection (Kicking)

Even though this is a very naïve algorithm (forcing GAD to take big steps), it had 40% convergence, much higher than before.

SCINE Results

40% Convergence

λ_0 : -3.49 ± 3.19 eV/ \AA^2

λ_1 : -0.21 ± 0.38 eV/ \AA^2

Steps: 1134.5 ± 471.4

Time: 9.92 ± 4.29 s

HIP Results

40% Convergence

λ_0 : -10.85 ± 4.49 eV/ \AA^2

λ_1 : -0.43 ± 1.01 eV/ \AA^2

Steps: 1092.0 ± 516.8

Time: 75.70 ± 35.91 s

Typical Failure Modes:

- Start at floor displacement, good movement until ~ 600 steps, then oscillates (~ 7 neg)
- Some graphs with ranges 12neg–4neg: Explodes at ~ 1300 steps

5 Higher-Order GAD (Multi-Mode Escape)

(Smarter kicks)

GAD only considers the smallest eigenvector. What if we explore higher orders?

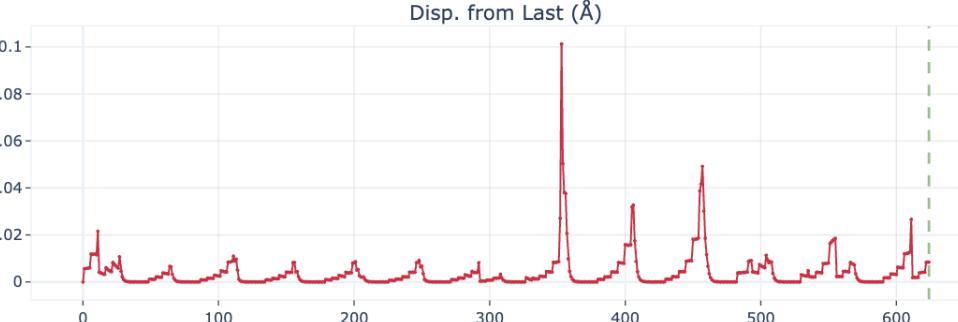
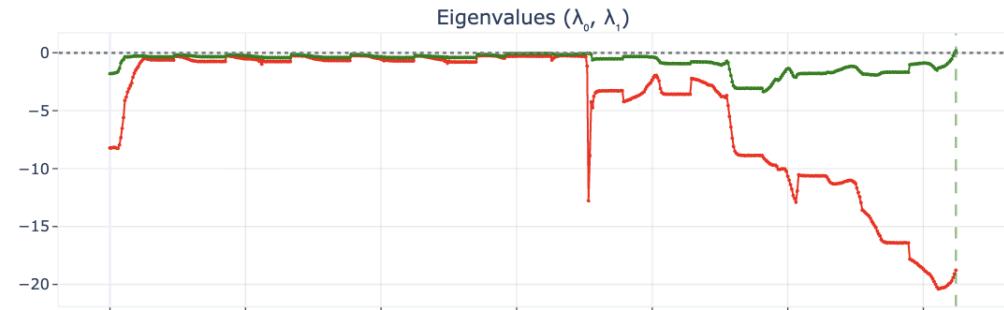
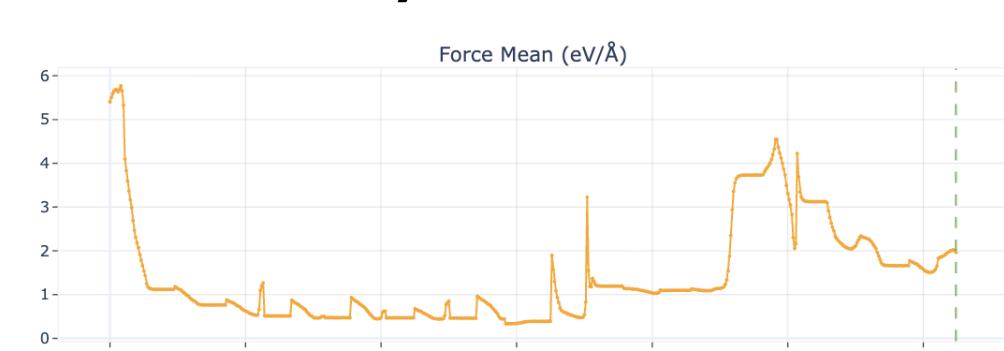
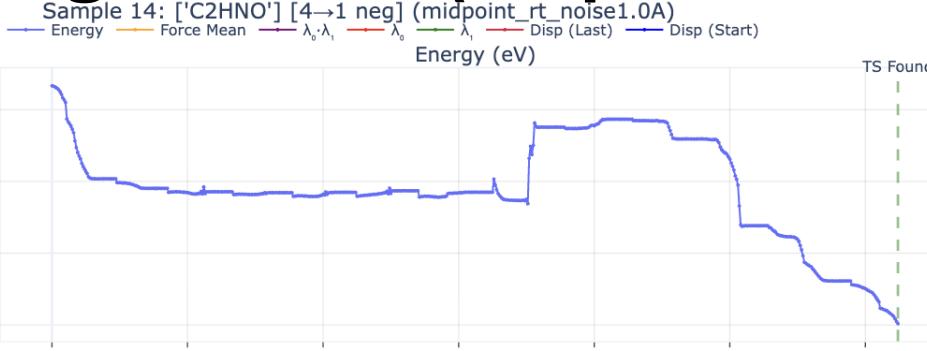
Key Insight

The “plateaus” didn’t make sense because the geometry was still unstable. **Discovery:** The eigenvalue being minimized was converging (force norm becoming tiny), but other negative modes remained.

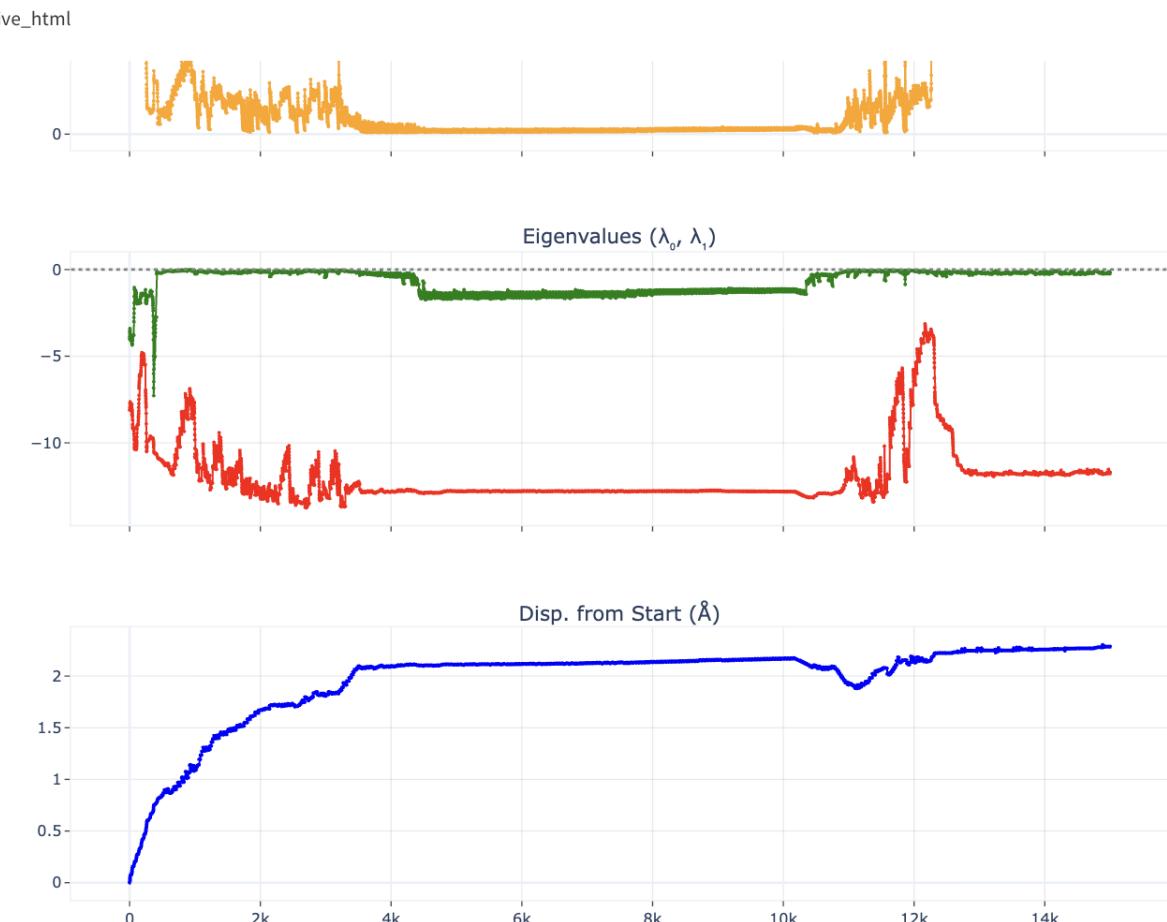
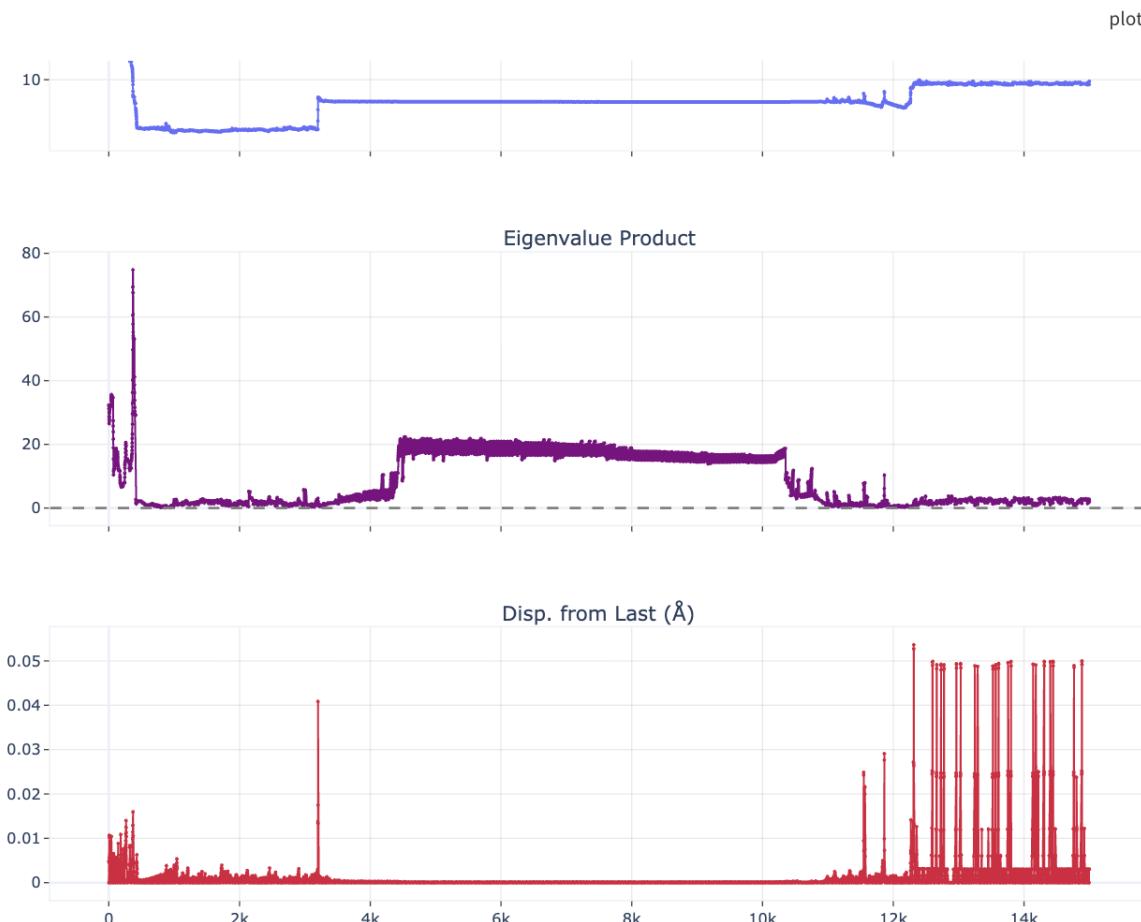
Solution: When stuck at a high-index saddle, perturb along v_2 (second eigenvector) to escape, then resume GAD.

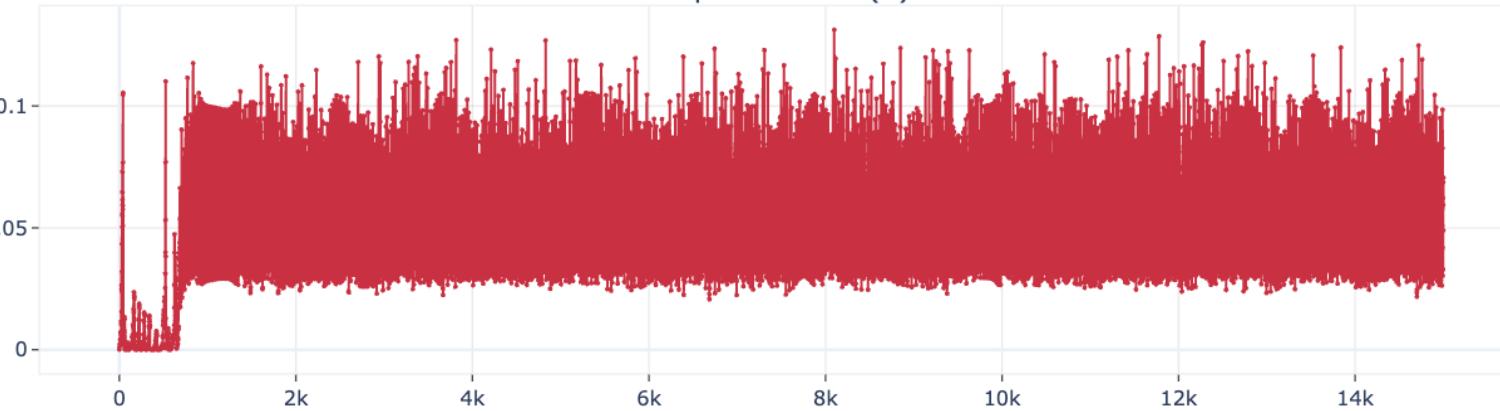
1. Run GAD until displacement becomes tiny (plateau detected)
2. If converged to order-1 saddle: success
3. If stuck at order > 1 : Perturb along v_2 (escape perturbation)
4. Resume GAD from perturbed geometry
5. Repeat until order-1 or max cycles reached

Look at the bottom left, we are constantly applying kicks as soon as the force direction becomes orthogonal to the smallest eigenvector (step size becomes too small).

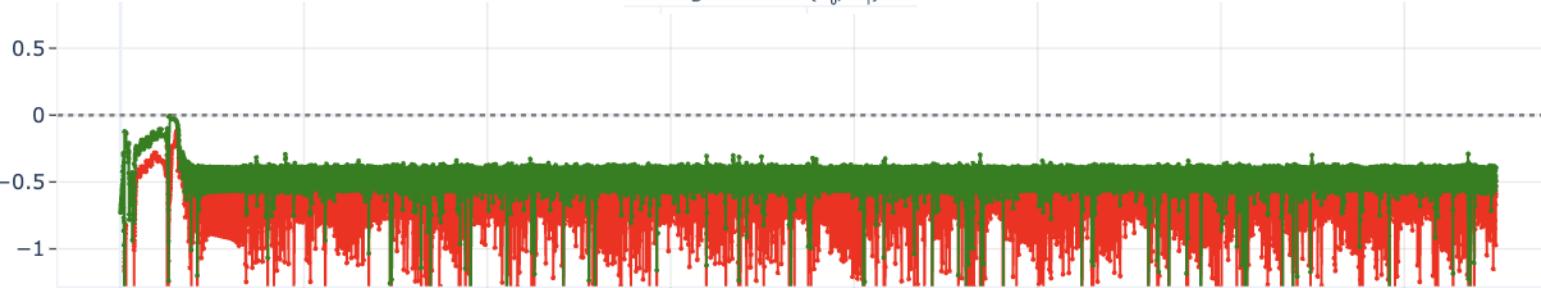
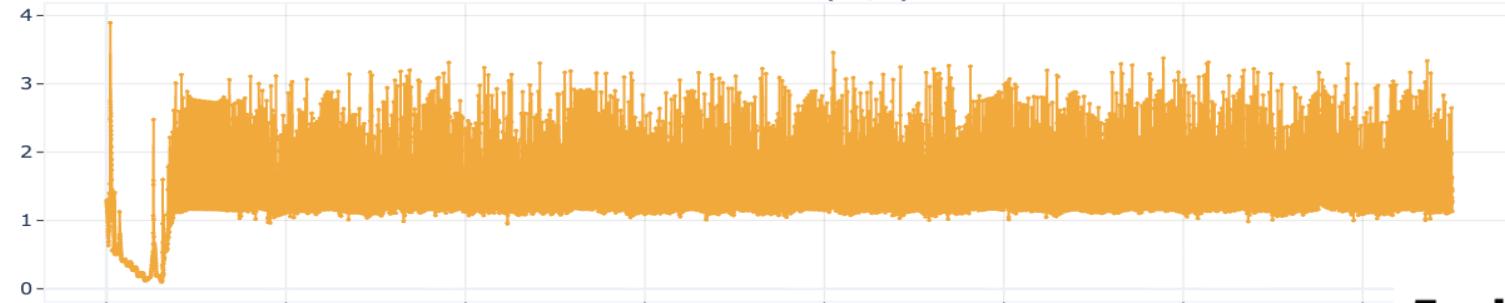


In this example, the trajectory just keeps oscillating. Hypothesize that this is because the “lowest eigenvalue” identity is being switched, because we’re not tracking modes.



Disp. from Last (\AA)

A closer look at the displacement, eigenvalues, and force mean that just keeps oscillating

Eigenvalues (λ_0, λ_1)Force Mean (eV/ \AA)

Even though we're still at 40% convergence, trajectories look much better (less plateaus, and getting much closer to convergence)

HIP Results

40% Convergence

Same comments as SCINE.

HIP Errors observed:

edge_vec_0_distance errors ($\sim 10^{-5}$)

λ_0 : -4.02 ± 4.18 eV/Å²

λ_1 : -0.71 ± 1.20 eV/Å²

Steps: 344.1 ± 175.6

Time: 24.19 ± 12.30 s

SCINE Results

40% Convergence

Same as kicking—but **different failure mode!**

Before: Moving in wrong directions (naive)

Now: 2–3 smallest eigenvalues oscillating with each other

λ_0 : -5.60 ± 9.81 eV/Å²

λ_1 : -0.25 ± 0.52 eV/Å²

Steps: 365.7 ± 82.3

Time: 2.84 ± 0.55 s

Two tweaks to try and get multi-mode GAD to work:

(based on observations from trajectories)

Key Insight

Two critical improvements that enable robust convergence:

- 1. Mode Tracking:** When eigenvalues are close or degenerate, eigenvector ordering can swap between steps, causing oscillations. Track the eigenvector with maximum overlap to previous step.
- 2. Trust Radius:** Limit maximum displacement per step to prevent “explosions” into unphysical regions.

Combined Effect:

- Mode tracking eliminates oscillations from eigenvector reordering
- Trust radius prevents geometry explosions after escape perturbations
- Together: stable traversal through high-dimensional saddle landscape

SCINE Results

92%

λ_0 : $-4.02 \pm 5.77 \text{ eV}/\text{\AA}^2$
 λ_1 : $-0.006 \pm 0.07 \text{ eV}/\text{\AA}^2$
Steps: 3934.9 ± 5032.0
Time: $35.67 \pm 45.53 \text{ s}$

HIP Results

71%

λ_0 : $-5.22 \pm 5.43 \text{ eV}/\text{\AA}^2$
 λ_1 : $-0.30 \pm 0.84 \text{ eV}/\text{\AA}^2$
Steps: 3858.8 ± 5636.1
Time: $305.75 \pm 390.02 \text{ s}$

Sample that was previously stuck in an infinite oscillation loop, now converging

